

Distributed Systems (CS 543)

Networks, RMI, & Pub/Sub

Dongman Lee

Dept of CS

KAIST

Class Overview

- Introduction
- Networking Issues for Distributed Systems
- Network Basics
- Remote Method Invocation
- Message-based Invocation: Publish/Subscribe

Introduction - Definition

- Definitions
 - A set of machines (computers, phones, routers, switches, gateways, etc.) connected by communication links
 - A communication subsystem
- Why networks?
 - Sharing resources
 - Increasing proximity between people

Introduction - Network Basics

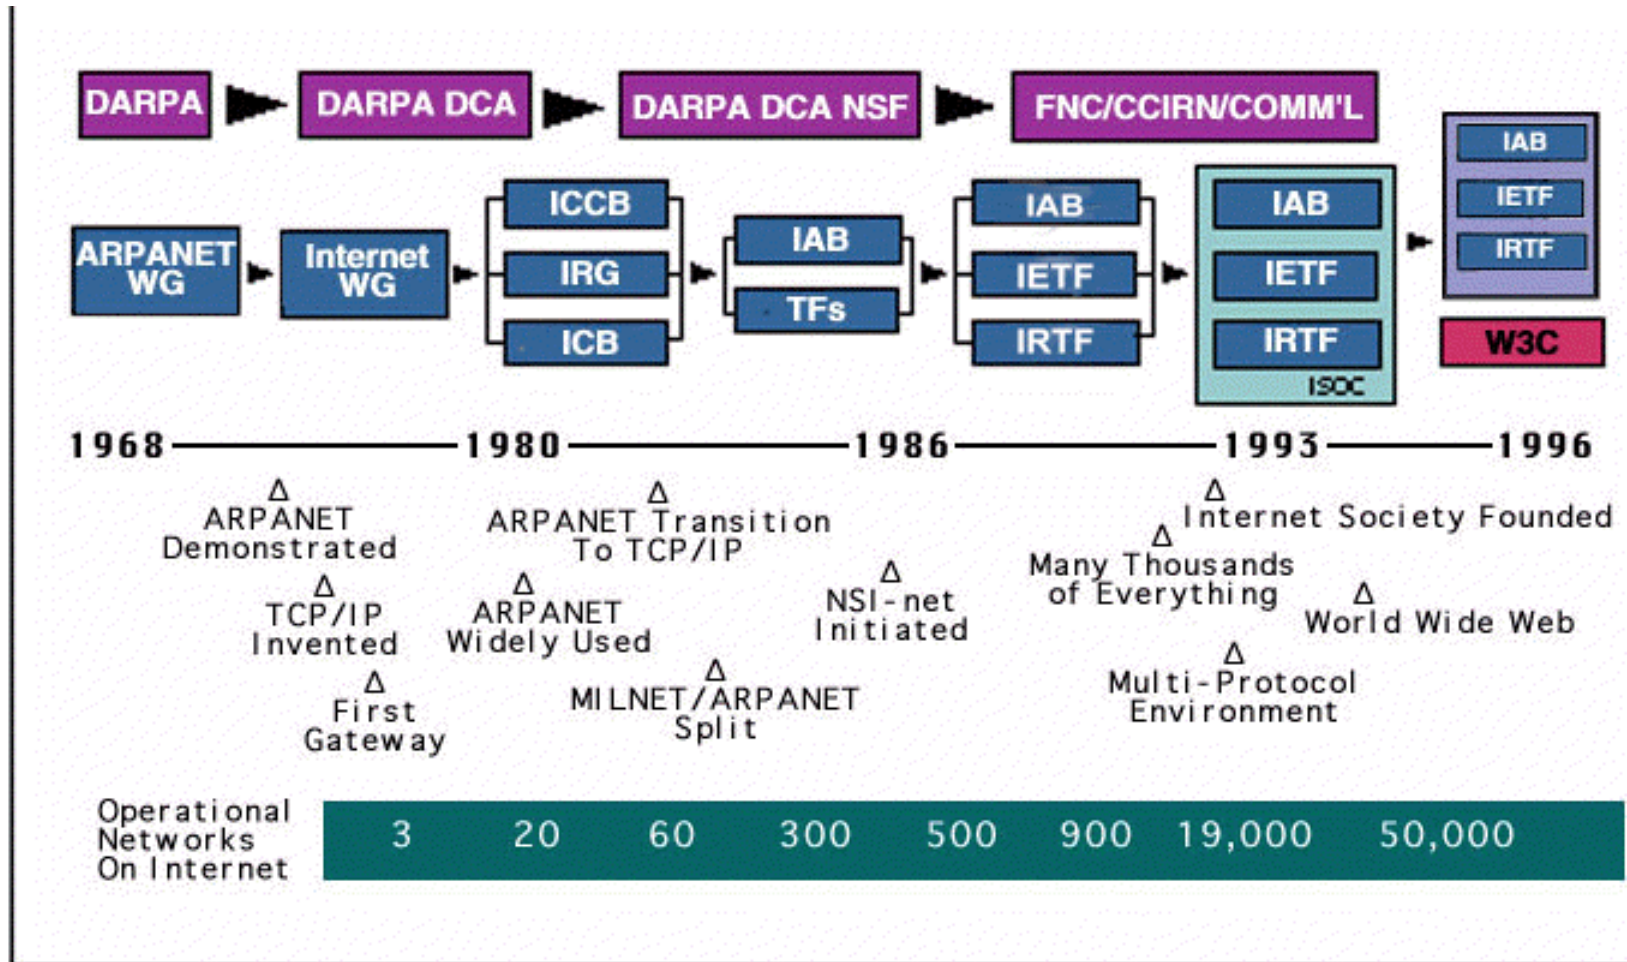
	<i>Example</i>	<i>Range</i>	<i>Bandwidth (Mbps)</i>	<i>Latency (ms)</i>
<i>Wired:</i>				
LAN	Ethernet	1-2 kms	10-1000	1-10
WAN	IP routing	worldwide	0.010-600	100-500
MAN	ATM	250 kms	1-150	10
Internetwork	Internet	worldwide	0.5-600	100-500
<i>Wireless:</i>				
WPAN	Bluetooth (802.15.1)	10 - 30m	0.5-2	5-20
WLAN	WiFi (IEEE 802.11)	0.15-1.5 km	2-54	5-20
WMAN	WiMAX (802.16)	550 km	1.5-20	5-20
WWAN	GSM, 3G phone nets	worldwide	0.01-02	100-500

Introduction - History of Networking

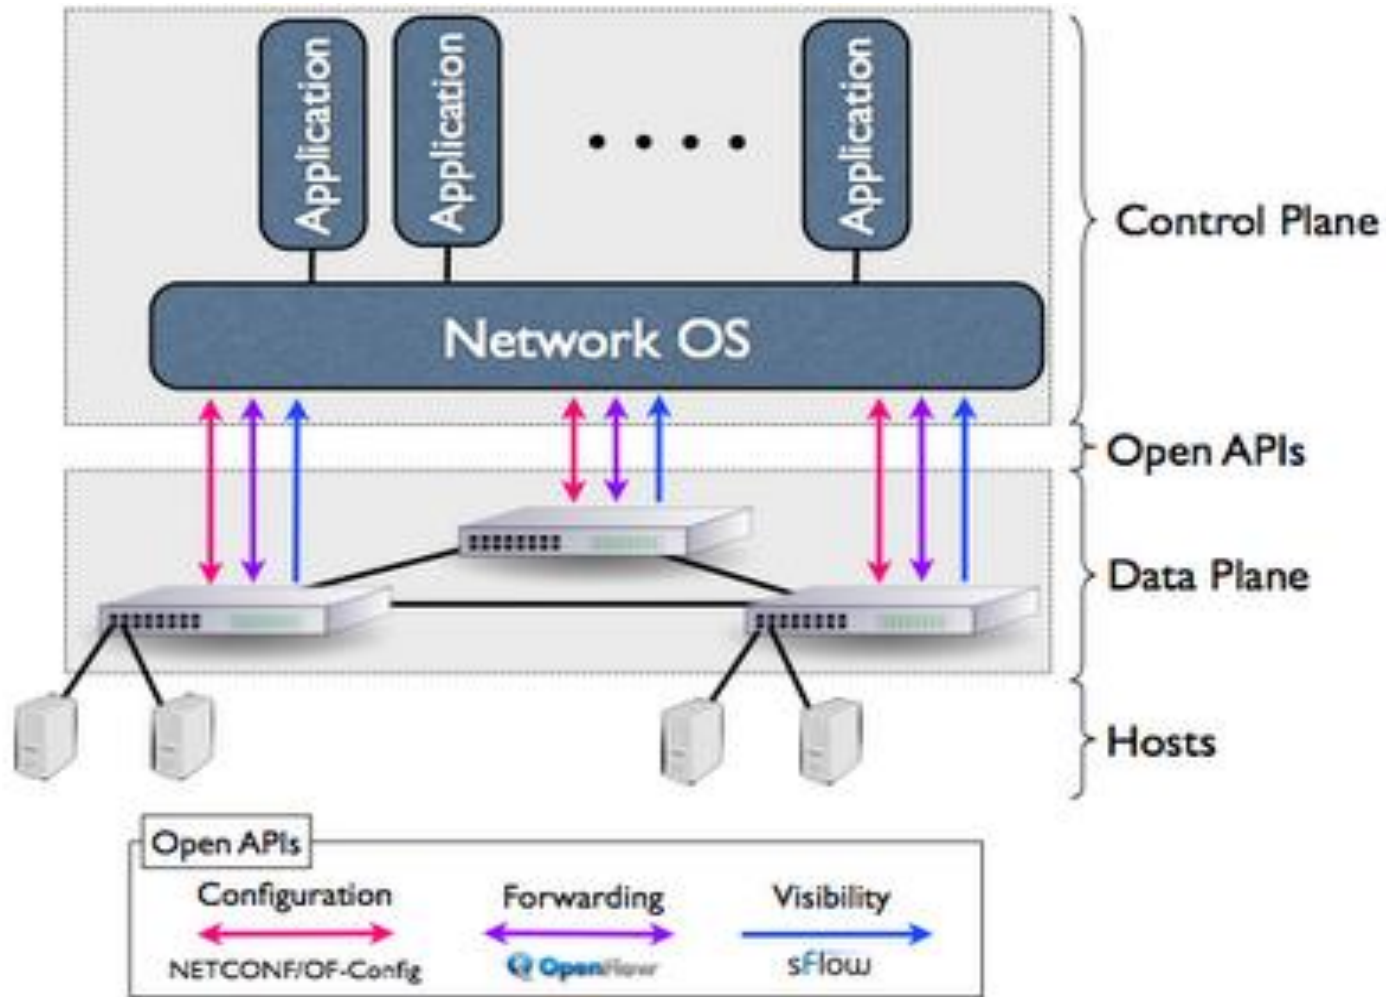
- Initially circuit switching (PSTN)
- 1960's: packet-switching
 - ARPAnet (1969)
- 1970's: local area network (LAN)
 - Ethernet
- 1980's: workstations & PCs
 - proliferation of LANs and WANs
- 1990's: WWW
 - explosion of Internet nodes; mobile networks
- 2000's: ad-hoc , sensor networks, PAN
 - Spontaneous; infra-less; energy efficiency
 - Programmable network: software defined network (SDN)

Introduction - History of Internet

- <http://www.isoc.org/internet/history>



Software Defined Networking



http://blog.sflow.com/2012_05_01_archive.html

Networking Issues for Distributed Systems

- Performance
 - transmission delay is arbitrary but finite
 - transmission time
 - ◆ latency + message size/ data transfer rate
 - total system bandwidth
 - ◆ LAN vs. WAN
 - local vs. remote operations
- Scalability
 - number of hosts and networks
 - ◆ addressing, routing mechanism, domain names
 - Latency
- Security
 - firewall
 - VPN

Networking Issues for Distributed Systems (cont.)

- Reliability
 - hardware medium does not support replication
 - error recovery
 - ◆ message lost
 - ◆ message duplication
 - ◆ message out-of-order
 - ◆ message corruption
- Quality of service
 - bandwidth, latency, jitter, and reliability
 - admission control and resource reservation
- Mobility
 - location management
 - hand-off
 - energy

Protocols

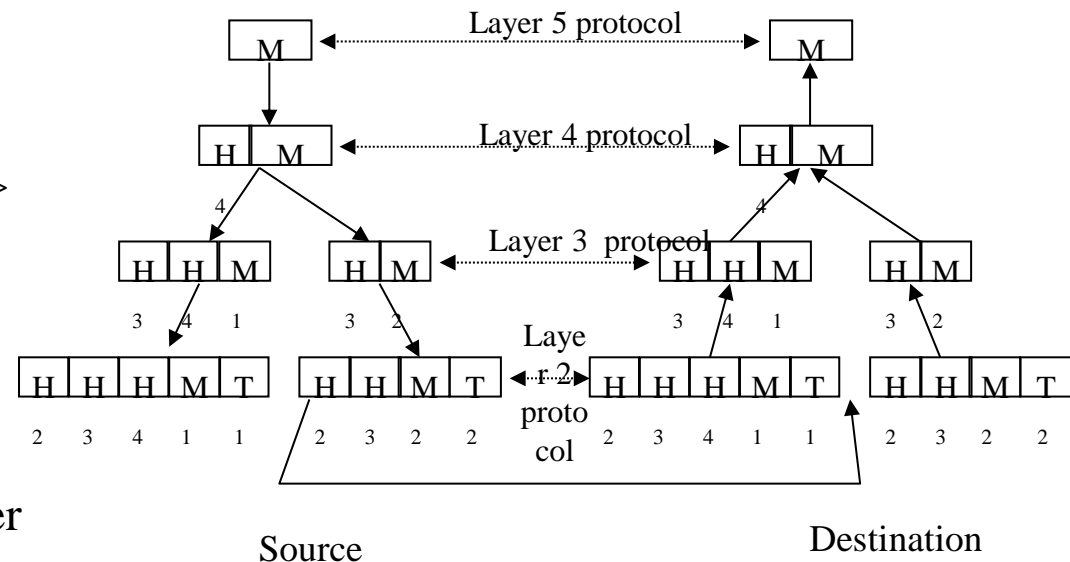
- A protocol specifies
 - rules and formats used for interactions between two parties
 - ◆ specification of the sequences of messages
 - ◆ specification of the format of the data in the message
- Requirements
 - specifications must be represented in finite states
 - implementation independent

Layered Architecture

- Motivation
 - reducing complexity by modularizing tasks vertically
 - encapsulation by a well-defined service interface to the upper layer
 - ◆ layer N provides a service to layer N+1
 - ◆ layer N extends the service of layer N-1
 - independence of each layer allows various implementations across layers
- Protocol suites
 - a complete set of protocol layers
 - protocol stack


Services in Layered Protocols

- Data transport service types
 - connection-oriented vs. connectionless
 - virtual circuit vs. datagram
 - depends on quality of service (QoS) requirements from the layer above
- Packet assembly
 - message at layer N+1 is fragmented into multiple PDUs at layer N (MTU at layer N+1 > MTU at layer N)
 - layer N at peer: responsible for assembly
- Addressing
 - identification of the peer at layer N
 - conversion into layer N-1 address



Abstraction of Communication Subsystem

- Regards a network as another I/O device
 - no abstraction beyond a transport protocol
 - ◆ client and server are wholly responsible for message exchange over a given transport mechanism (e.g. TCP/IP)
 - ◆ primitives: write (send) and read (receive)
- Request and reply
 - abstraction of message passing required to execute a procedure at a server
 - ◆ primitives: DoOperation, GetRequest, and SendReply
- Remote method invocation
 - hides the separation between a client and a server
 - ◆ makes the invocation of a remote method residing at a server the same as that of a local one

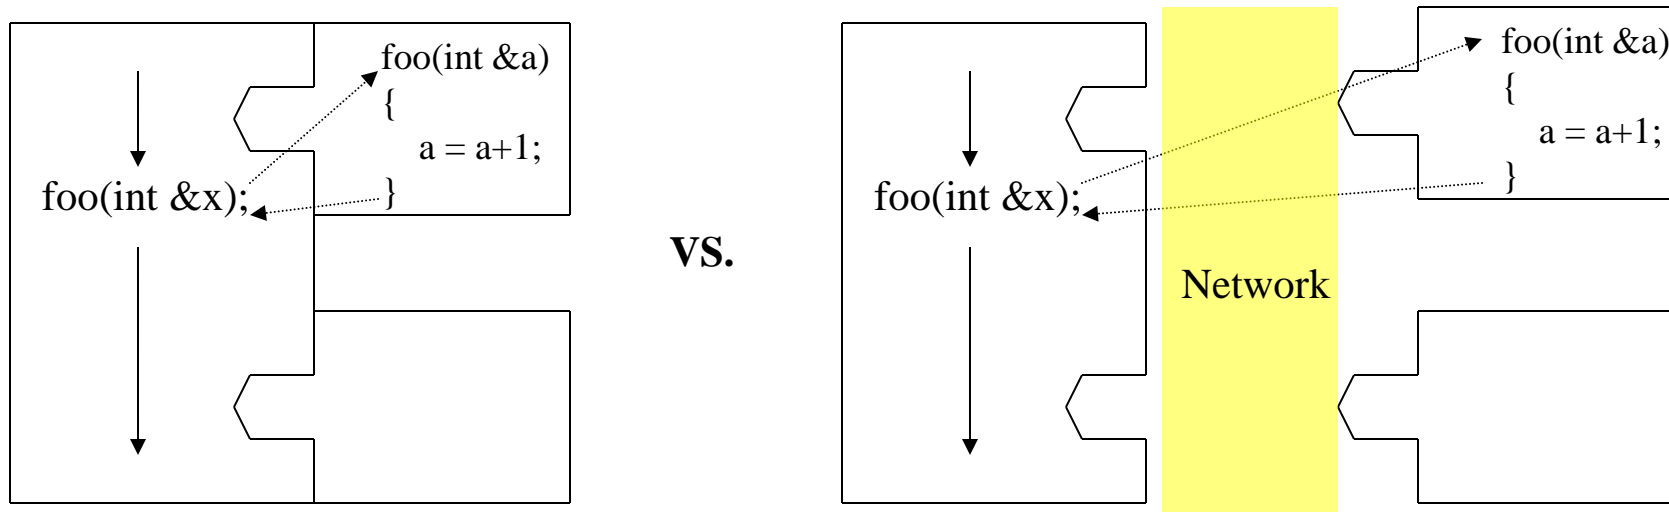


RMI- Motivation

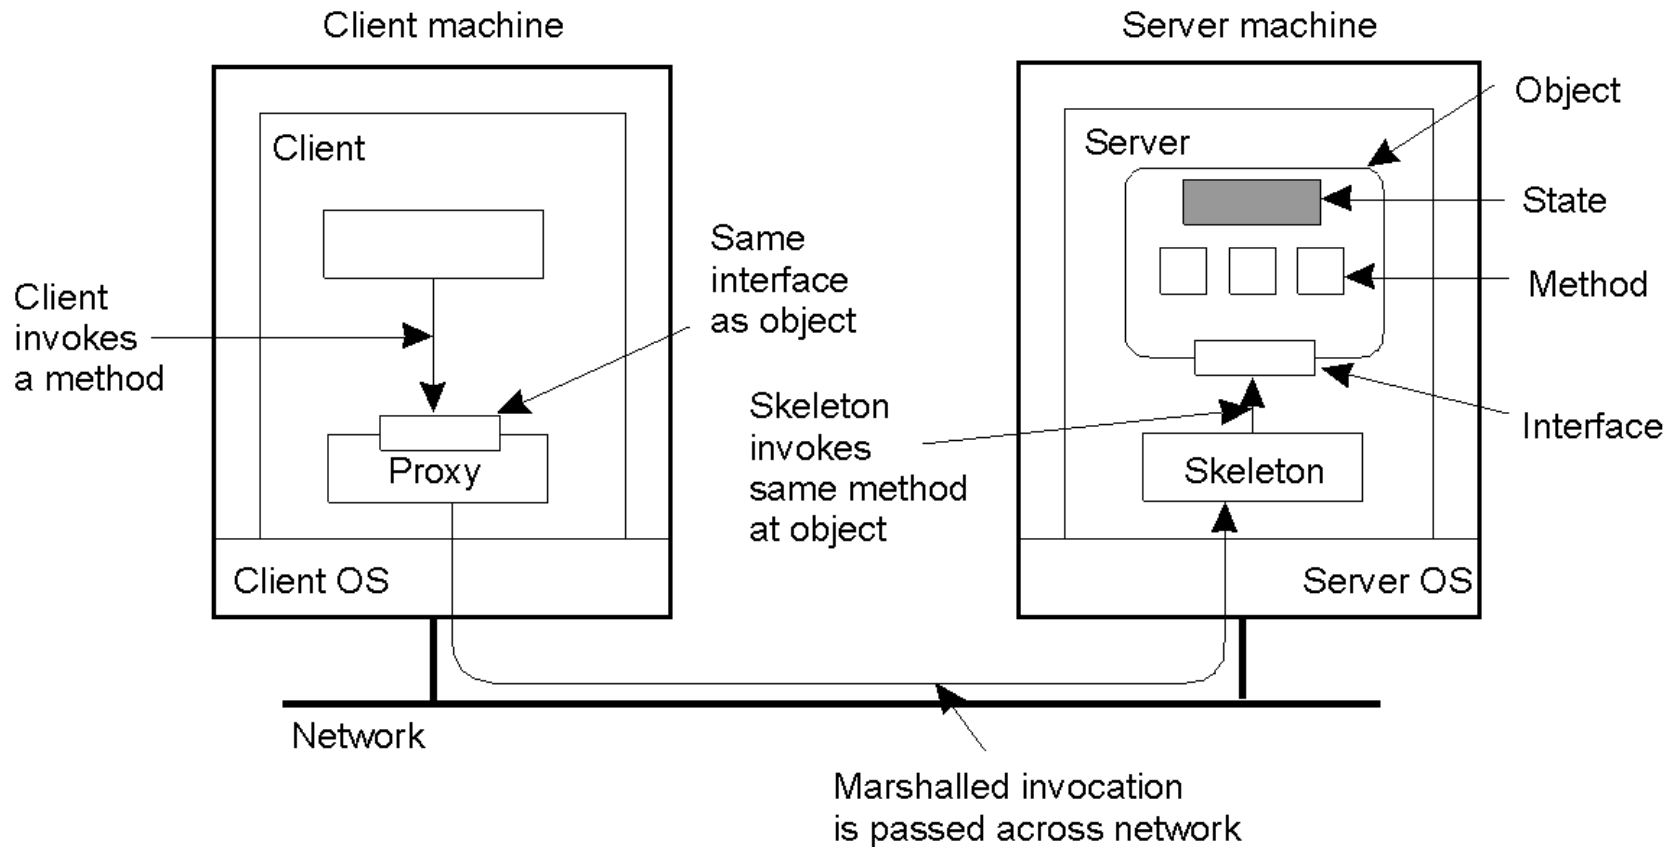
- Construction of distributed applications is a difficult task
 - Heterogeneity
 - Separation
 - Partial failure
- Leverage a local method invocation [Birrell & Nelson]
 - method invocations are a well-known and well-understood mechanism for transfer of control and data within a program
 - the same principle can be extended to the client-server interaction
 - ◆ clean and simple semantics
 - ◆ efficiency
 - ◆ generality

Local Invocation vs. Remote Invocation

- Client and server fail independently
 - extra exception handling
 - return value needs to be overloaded with failure codes from the server
- Global variables and pointers cannot be used across the interface



Remote Invocation Structure



Distributed object method invocation

Design Considerations

- Abstraction of RPC semantics
- Heterogeneity
- Parameter passing
- Binding
- Failure Transparency
- Support for concurrency

Level of Abstraction of RMI Semantics

- Integrate into the language
 - extend language to provide constructs for remote invocation semantics
 - examples: Cedar, Argus, Mercury, Java RMI
- Describe in a special purpose interface definition language
 - interface definition language supports language independence
 - ◆ limited representation of parameter types
 - examples:
 - ◆ Sun XDR, DCE IDL, Xerox Courier,
 - ◆ CORBA IDL, DCOM MIDL

Level of Abstraction of RMI Semantics (cont.)

- IDL provides a distributed object with an abstract way to identify
 - what the server interface is going to be
 - what the interface will support
 - what other interface(s) it will collaborate with
- A client does not need to know how the interface is implemented
 - a client only incorporates the stub code generated from the IDL specification in the client implementation language
 - a client does not get any impact from change(s) in the server implementation as long as the server interface stays the same
 - the server can have multiple implementation of the same interface

Example: CORBA IDL

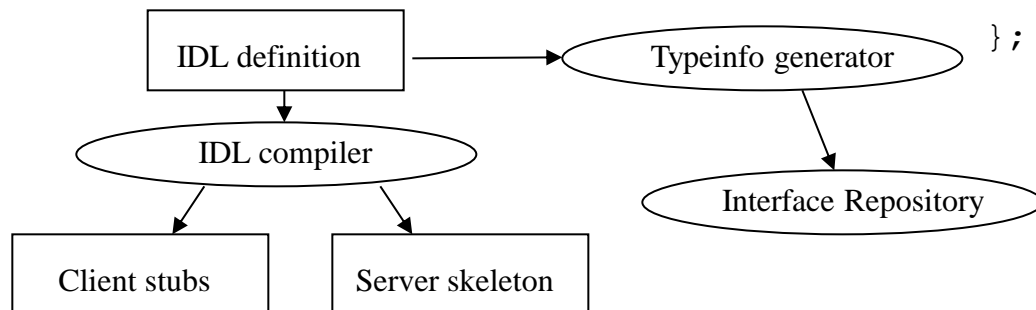
- IDL is the language used to describe the interfaces that client objects can call and object implementations provide

- C++-like syntax
- features
 - ◆ multiple inheritance
 - ◆ attributes
 - ◆ exceptions
 - ◆ contexts

- CORBA has two IDL compilers

- stub & skeleton generator
- typeinfo generator for dynamic invocation

```
module Shopping {  
    /* structure of an item */  
    struct Item {  
        string item_name;  
        float  item_price;  
    };  
    typedef sequence <Item> Basket;  
    /* class GoShopping */  
    interface GoShopping {  
        exception NotEnoughMoney  
            {float  
needed};  
        attribute float money;  
        void buy (in Basket, out Basket)  
            raises  
(NotEnoughMoney);  
    };  
};
```



Heterogeneity

- Marshalling policy
 - conversion to the common data type agreed by both a sender and a receiver: Sun RPC, CORBA
 - a receiver makes it right (tagging): NCS RPC
 - a sender negotiates with a receiver

3 0	2 0	1 0	0 5
7 L	6 L	5 I	4 J

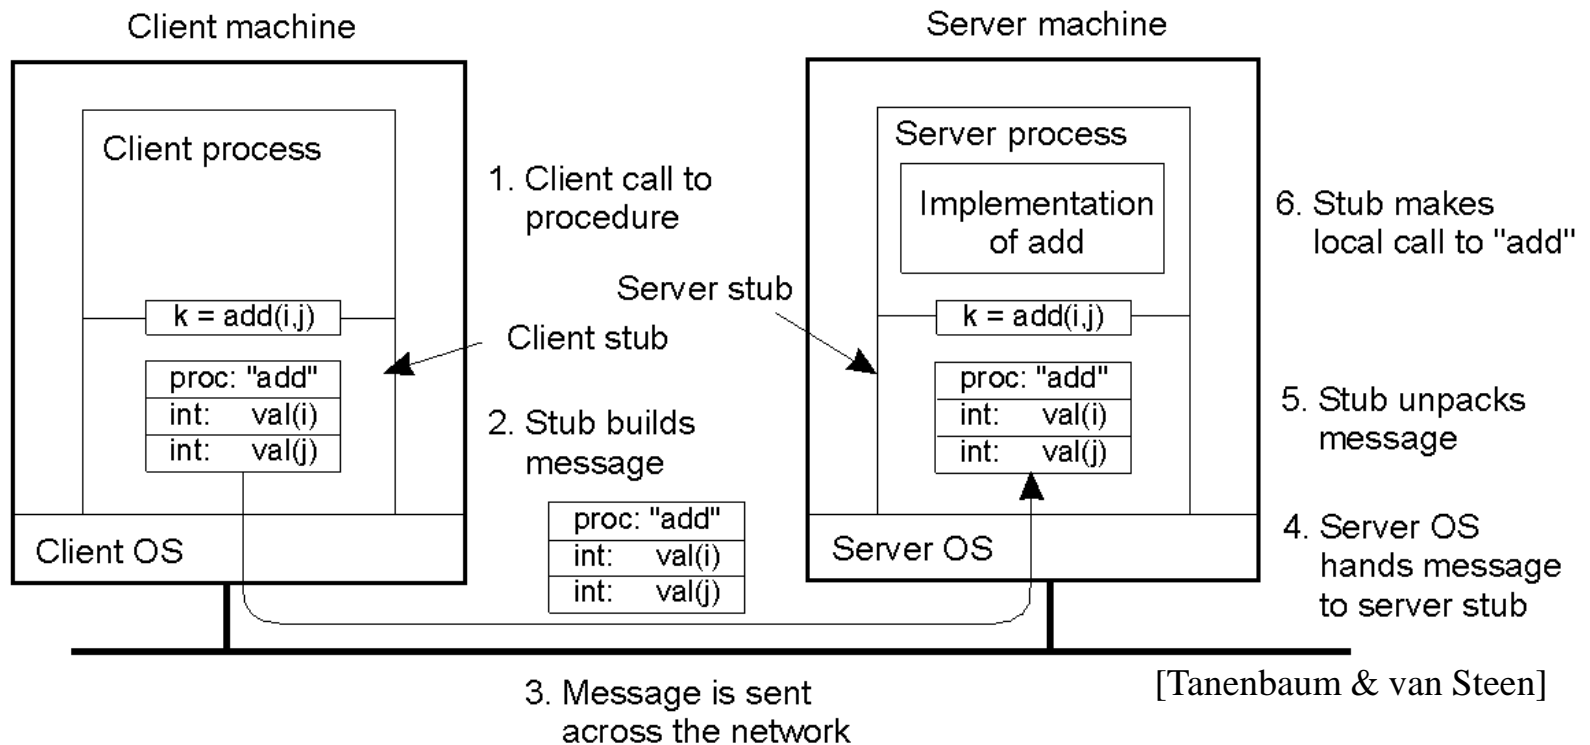
0 5	1 0	2 0	3 0
4 J	5 I	6 L	7 L

- Accommodate multiple transport protocols

Parameter Passing

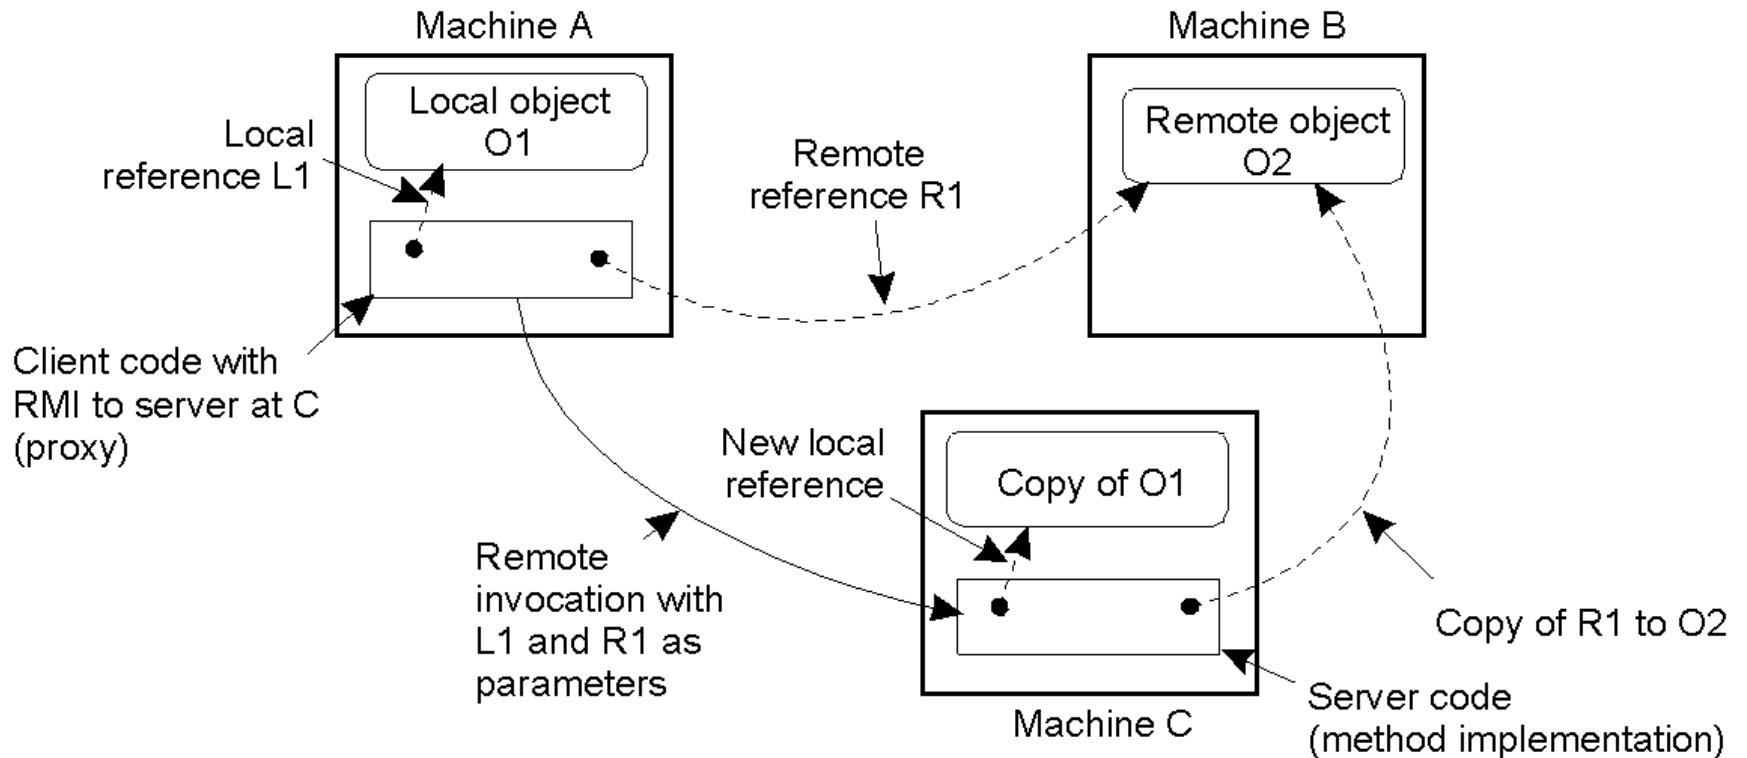
- Considerations for parameter passing in RMI
 - direction of parameters
 - ◆ input: client -> server
 - ◆ output: server -> client
 - ◆ input & output: client <-> server
 - call-by-reference
 - ◆ use call-by-value-result instead
 - pointer parameters
 - ◆ no meaning to pass a pointer to memory location
 - ◆ linearize before passing it
 - procedure parameter
 - ◆ procedure should be accessible across the network

Parameter Passing (cont.)



Parameter Passing (cont.)

- Distributed objects




Binding

- Naming
 - mapping from a name to a particular service
 - ◆ static linking: most RPCs
 - ◆ dynamic linking: NCS RPC
 - ◆ procedure variable: Argus, Mercury
 - exportation and importation
 - ◆ server exports its service interface
 - ◆ client imports exported service
 - interface consists of
 - ◆ name: uniquely identifiable
 - version info and location info can be included
 - ◆ signature: parameter name & type

Binding (cont.)

- Locating a server
 - via a binder (like a directory service)
 - ◆ well-know address
 - ◆ run-time notification
 - ◆ broadcasting (recruiting)



Failure Transparency


- Ordered delivery
 - request and reply delivered in order (e.g. Argus)
- RPC Execution semantics
 - How many times will a remote procedure be executed by the callee per call?
 - Semantics
 - ◆ at-least-once; at-most-once; exactly-once
- At least once
 - when a caller receives the expected result, it implies that the requested call has been executed one or more (at least once) at a callee
 - the callee does not remember multiple executions due to call duplication or retry after failure
 - acceptable if operations are idempotent
 - example: Sun RPC – NFS

Failure Transparency (cont.)

- At-most-once
 - when a caller receives the expected result, it implies that the requested call has been executed exactly once. If not completed, executed zero or one time
 - the callee can detect duplicated requests at the absence of failure but persist through the abnormal termination (failure amnesia)
 - example: NCS RPC
- Exactly-once
 - when a caller receives the expected result, it implies that the requested call has been executed exactly once. If not completed, executed zero times
 - the callee is failure-resilient: logging and two-phase commit
 - example: Argus

Support for Concurrency

- RMI behaves differently from LMI
 - the client does not need to be blocked while the server executes the requested call
 - the client may not need the result
- Approaches
 - Synchronous
 - ◆ a process/thread per call (at a client or a server)
 - Asynchronous
 - ◆ two types
 - no result
 - delayed result: delayed synchronous
- Examples:
 - Argus, Mercury(Promises), X11 RPC, ANSA REX
 - CORBA
 - ◆ oneway



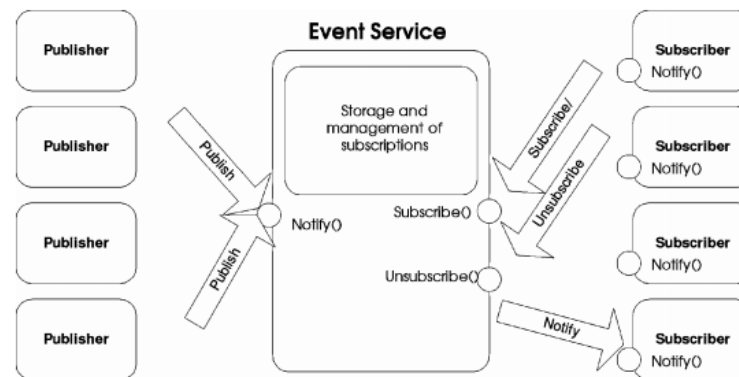
Why Publish/Subscribe?

- Demand for more flexible communication models
 - Individual *point-to-point* and *synchronous* communications lead to rigid and static applications in large-scale distributed systems
 - Increasing Scale of distributed systems
 - ◆ The Internet has considerably increased the scale of distributed systems
 - Dynamicity
 - ◆ Entities in distributed systems vary their location and behavior throughout the lifetime.



Publish/Subscribe

- Loosely coupled interaction scheme
 - Subscribers
 - ◆ Subscribers express their interest in any events, and are subsequently notified of matched events, generated by a publisher
 - Publishers
 - ◆ Publishers generate and publish new information denoted by an event
 - Event notification services
 - ◆ A neutral mediator between publishers and subscribers which provides storage and management for subscriptions and efficient delivery of events
 - Events
 - ◆ An event from a publisher is asynchronously propagated to all subscribers that registered interest in that given event



Event Service Decoupling


- Event service decoupling
 - Event services provide the full decoupling between publishers and subscribers
 - The full decoupling releases the burden of application designers in large-scale systems
 - The strength of publish/subscribe lies in decoupling:
 - ♦ Decoupling increases scalability by removing all explicit dependencies between the interacting participants
 - ♦ Dependency-free infrastructure can well adapted to distributed environments that are asynchronous by nature (e.g. Mobile environment)
 - Event service decoupling can be decomposed along the following three dimensions:
 - ♦ Time
 - ♦ Space
 - ♦ Synchronization

Event Service Decoupling (cont.)


- Decoupling dimensions
 - Space
 - ◆ The interacting parties do not need to know each other's references or number
 - Time
 - ◆ The interacting parties do not need to be actively participating in the interaction at the same time
 - ◆ Publishers can publish events while the subscribers are disconnected, and vice versa
 - Synchronization
 - ◆ Publishers are not blocked while producing events
 - ◆ Subscribers can get asynchronously notified of the occurrence of an event through a callback



Space decoupling



Time decoupling



Synchronization decoupling

Publish/Subscribe Variations

- Publish/Subscribe Variations
 - Subscribers are usually interested in particular events or event patterns, and not in all events
 - The different ways of specifying the events of interest have led to several subscription schemes:
 - ◆ Topic-based publish/subscribe
 - ◆ Content-based publish/subscribe
 - ◆ Type-based publish/subscribe

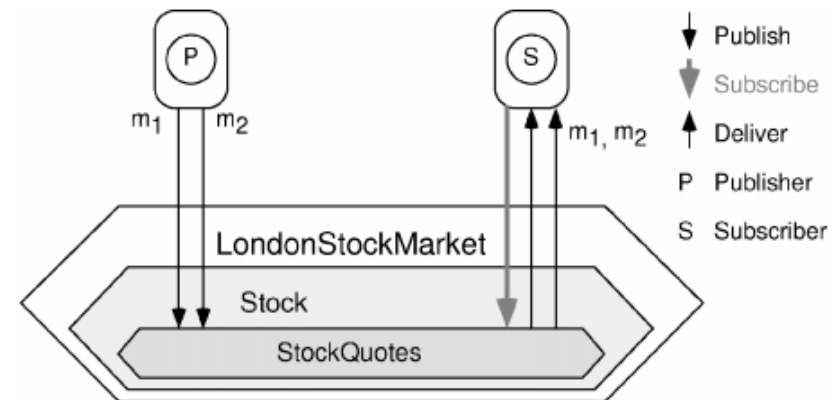
Publish/Subscribe Variations (Cont.)

- Topic-based Publish/Subscribe (*or Subject-based*)
 - The earliest publish/subscribe scheme which extends the notion of *channels*
 - Subscriptions
 - ♦ Participants can publish events and subscribe to individual topics, which are identified by *keywords*
 - ♦ Every topic is viewed as an event service of its own, identified by a unique name, keyword
 - Similarity to *group communications*
 - ♦ Subscribing to a topic $T \approx$ Becoming a member of a group T
 - ♦ Publishing an event on topic $T \approx$ Broadcasting that event among the members of T
 - Improvement: *Hierarchies*
 - ♦ Topics can be organized according to containment relationships (*hierarchical addressing*)
 - ♦ A subscription made to a node involves those to all the subtopics

Publish/Subscribe Variations (Cont.)

- Topic-based Publish/Subscribe (*or Subject-based*)
 - Sample code and an interaction example

```
public class StockQuote implements Serializable {  
    public String id, company, trader;  
    public float price;  
    public int amount;  
}  
public class StockQuoteSubscriber implements Subscriber {  
    public void notify(Object o) {  
        if (((StockQuote)o).company == 'TELCO' && ((StockQuote)o).price < 100)  
            buy();  
    }  
}  
// ...  
Topic quotes = EventService.connect("/LondonStockMarket/Stock/StockQuotes");  
Subscriber sub = new StockQuoteSubscriber();  
quotes.subscribe(sub);
```



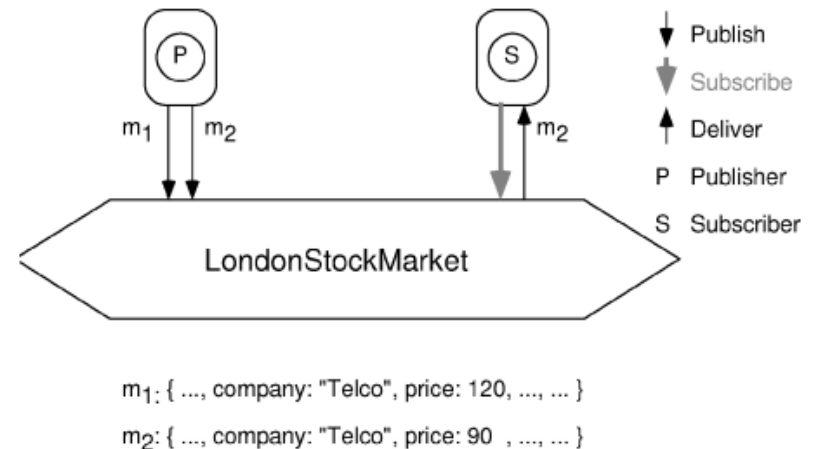
Publish/Subscribe Variations (Cont.)

- Content-based publish/subscribe (*or property-based*)
 - Topic-based scheme offers only limited expressiveness
 - Subscriptions
 - ◆ Subscribers subscribe to the actual content of the considered events
 - ◆ Events are not classified according to some predefined external criterion (e.g. topic name), but according to the properties of the events themselves
 - Selective subscriptions with *filters*
 - ◆ Subscribers specify filters using a subscription language
 - ◆ Filters define constraints, usually in the form of name-value pairs of properties and basic comparison operators (e.g., =, <, >, etc.)
 - ◆ Constraints can be logically combined (e.g., AND, OR)

Publish/Subscribe Variations (Cont.)

- Content-based publish/subscribe (*or property-based*)
 - Sample code and an interaction example

```
public class StockQuote implements Serializable {  
    public String id, company, trader;  
    public float price;  
    public int amount;  
}  
public class StockQuoteSubscriber implements Subscriber {  
    public void notify(Object o) {  
        buy();    // company == 'TELCO' and price < 100  
    }  
}  
// ...  
String criteria = ("company == 'TELCO' and price < 100");  
Subscriber sub = new StockQuoteSubscriber();  
EventService.subscribe(sub, criteria);
```



Publish/Subscribe Variations (Cont.)

- Type-based publish/subscribe
 - Event type
 - ◆ Topics usually regroup events that present commonalities not only in content, but also in structure
 - ◆ A regrouped events which share the commonalities are said to be of the same *type*
 - Advantages
 - ◆ This enables a *closer integration* of the language and the middleware (e.g., an object written in a language has its own structure as well as contents)
 - ◆ *Type safety* can be ensured at compile-time by parameterizing the resulting abstraction interface by the type of the corresponding events

Publish/Subscribe Variations (Cont.)

- Type-based publish/subscribe
 - Sample code and an interaction example

```
public class LondonStockMarket implements Serializable {  
    public String getId() {...}  
}  
public class Stock extends LondonStockMarket {  
    public String getCompany() {...}  
    public String getTrader() {...}  
    public int getAmount() {...}  
}  
public class StockQuote extends Stock {  
    public float getPrice() {...}  
}  
public class StockRequest extends Stock {  
    public float getMinPrice() {...}  
    public float getMaxPrice() {...}  
}  
public class StockSubscriber implements Subscriber<StockQuote> {  
    public void notify(StockQuote s) {  
        if (s.getCompany() == 'TELCO' && s.getPrice() < 100)  
            buy();  
    }  
}  
// ...  
Subscriber<StockQuote> sub = new StockSubscriber();  
EventService.subscribe<StockQuote>(sub);
```

Type hierarchy:

