

Distributed Systems

Naming

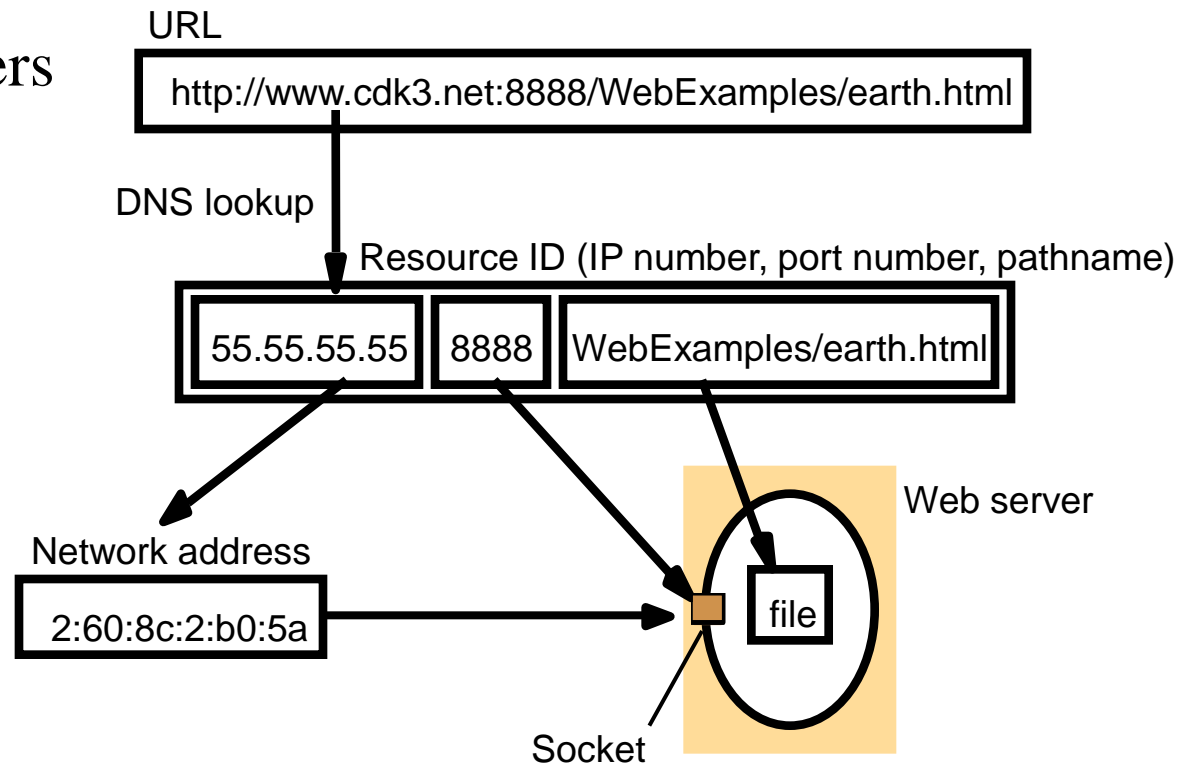
Dongman Lee
Dept of CS
KAIST

Class Overview

- Why Naming?
- Naming Fundamentals
- Name Services
- Directory & Discovery Services
- Locating Mobile Entities
- Case Studies
 - Host naming: DNS
 - Content naming: Naming in ICN
 - Directory: X.500
 - Discovery: SLP

Names are

- Network addresses
- Process identifiers
- File names
- User names
- Service names
- Class identifiers
- ...



Why Naming?

- Names are used to
 - facilitate sharing
 - ◆ a name is a common way for other objects to communicate with an object with that name
 - ◆ a single name may represent a set of multiple objects doing the same job
 - uniquely identify particular members of some large set
 - ◆ a name may specify roles played by the object(s) named
 - ◆ a name may give hints about the location of the object named
 - ◆ a name may indicate a set of operations allowed (authorized) by the object named

Naming in Distributed Systems

- Characteristics of naming in distributed systems
 - locations
 - ◆ are unintuitive
 - ◆ make object migration/mobility difficult
 - naming system could be distributed as well
 - ◆ its efficiency and scalability may depend on how distribution is done
- What we can learn from telephone networks [Birrel]
 - decentralized
 - changes are done by authorized entities
 - database is replicated
 - multiple names are allowed
 - names are context-dependent

Naming Fundamentals


- Different types of names
 - addresses
 - ◆ name of an access point to an associated entity
 - ◆ an entity may be accessible at more than one address
 - ◆ an entity may change its address and the address is reassigned to other entity
 - identifiers
 - ◆ properties of identifier
 - an identifier refers to at most one entity
 - an entity is referred to by at most one identifier
 - an identifier always refers to the same entity (immutable)
 - human readable names
 - ◆ represented as character-string
 - ◆ used in many different forms
 - DNS
 - file names

Naming Fundamentals (cont.)

- Name spaces
 - a collection of all valid names recognized by a name service
 - naming domain
 - ◆ a name space for which there exists a single overall naming authority for managing names within it
 - ◆ e.g. DNS TLD
 - naming context
 - ◆ an abstraction that either maps a given name to attributes or further naming context and derived name
 - ◆ e.g. CWD in UNIX
 - naming convention
 - ◆ absolute vs. relative
 - ◆ flat vs. hierarchical

Naming Fundamentals (cont.)

- Name spaces (cont.)
 - represented as a labeled, directed graph
 - ◆ regular node
 - represents a named entity and has no outgoing edges
 - contains information on the entity (address or identifier) or content of an entity
 - ◆ context node
 - has a number of outgoing edges with labels
 - contains context table of (edge label, node id) pairs
 - path name
 - ◆ a sequence of labels refers to a node
 - ◆ absolute path name
 - ◆ relative path name



Naming Fundamentals (cont.)

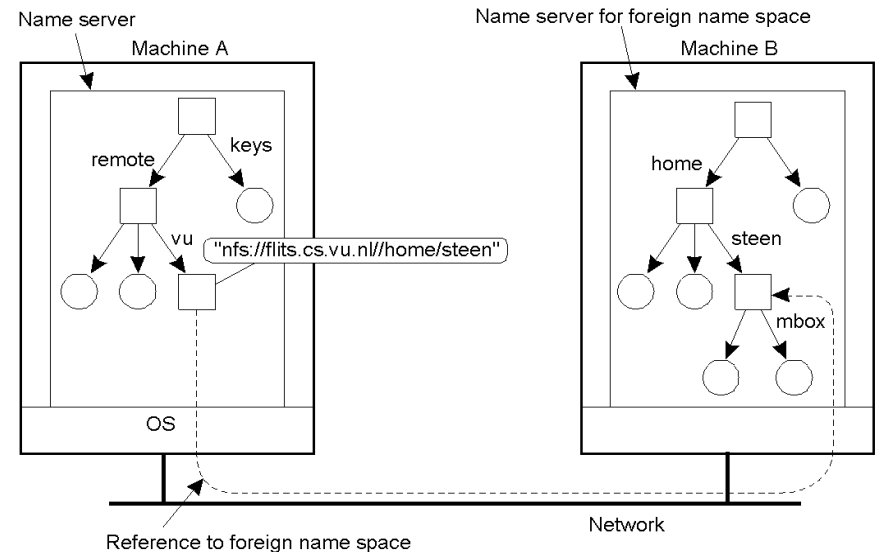
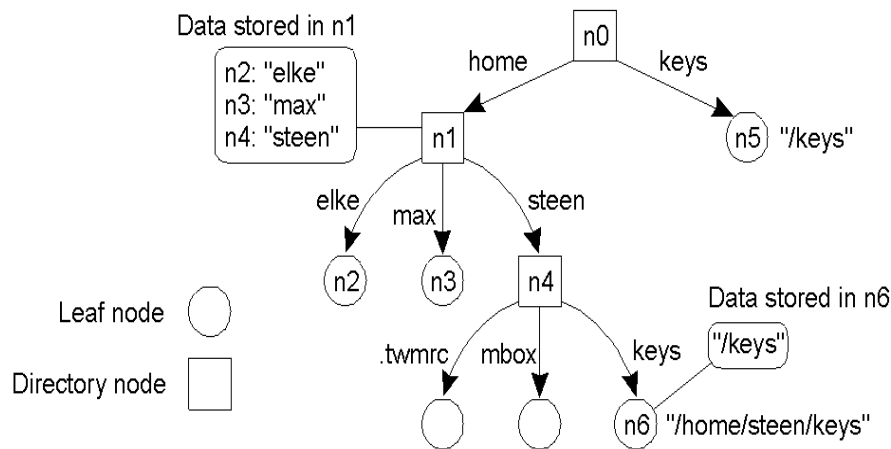
- Names and location
 - Non-pure names
 - ◆ names convey information on where to find the objects named
 - object location
 - name service location
 - Pure names
 - ◆ no interpretation of names themselves
 - ◆ pure names don't convey location

Naming Fundamentals (cont.)

- Name resolution
 - a process that looks up a name in the name database to obtain associated attributes
 - ◆ an object identifier or address
 - ◆ services provided by the object named
 - ◆ a server managing the object named
 - closure mechanism is required
 - ◆ knowing how and where to start name resolution
 - ◆ select the initial node in a name space to figure out where to start
 - system should provide a way of where to start
 - examples
 - » address of a root node of a file system
 - » pre-defined syntax in naming (e.g. phone numbers)
 - » name context hints (e.g. CWD, HOME)

Naming Fundamentals (cont.)

- Name resolution (cont.)
 - multiple levels of navigation
 - ◆ linking - aliases to another name contexts (e.g. symbolic or hard links)
 - ◆ mounting - merged name spaces (e.g. NFS or GNS)



Name Services

- Name service
 - a service that allows users and processes to add, remove, and look up names
 - main tasks
 - ◆ name space/database management
 - ◆ name resolution
- Desirable properties [Lampson]
 - scalable
 - long lifetime
 - high availability
 - fault isolation
 - tolerance of mistrust

Name Services (cont.)


- Design considerations [Needham]
 - bindings
 - ◆ relax binding time since distributed systems are more dynamic
 - ◆ cache frequently used bindings if possible and update on use
 - name resolution
 - ◆ name servers are usually replicated for availability and reliability
 - ◆ pure names should be used carefully considering the size of networks, broadcasting capability, etc.
 - consistency
 - ◆ accessibility is considered more important than absolute correct result
-> asynchronous consistency
 - scalability
 - ◆ never assume that the size of the name space is liable to get

Name Services (cont.)

- Name space management
 - name space can be partitioned into three levels [Cheriton/Mann]
 - ◆ global level
 - top level nodes in the name space
 - » context node encompasses many organizations
 - rare changes and large number of users
 - » cacheable in the lower levels - performance may be not relatively critical
 - » stability and high availability are important - replication
 - ◆ administrative level
 - context nodes representing group of nodes belonging to same organization in the name space
 - relatively stable but more frequent changes than global level nodes
 - » availability and performance requirements similar to global level
 - ◆ managerial level
 - context nodes for small group and regular nodes in the name space
 - frequent changes and high access rate
 - » lesser availability requirement but immediate response required


Name Services (cont.)

- Name resolution mechanisms
 - name resolver
 - ◆ agent that is responsible for ensuring name resolution process
 - iterative name resolution vs. recursive name resolution
 - Iterative name resolution
 - ◆ client's name resolver contacts as many name servers as it can until name resolution is done
 - ◆ each name server has partial information
 - ◆ intermediate server may play a role of name resolver



Name Services (cont.)

- Name resolution mechanisms (cont.)
 - recursive name resolution
 - ◆ name server passes the result to next name server as it can until name resolution is done
 - may put higher performance demand on each name server
 - ◆ allow each name server to gradually learn the address of other name servers responsible for part of names that it was not able to resolve
 - cached resolution results can be leveraged next time
 - ◆ may provide cheaper communication cost
 - administrative nodes are usually closer to their managerial nodes



Name Services (cont.)

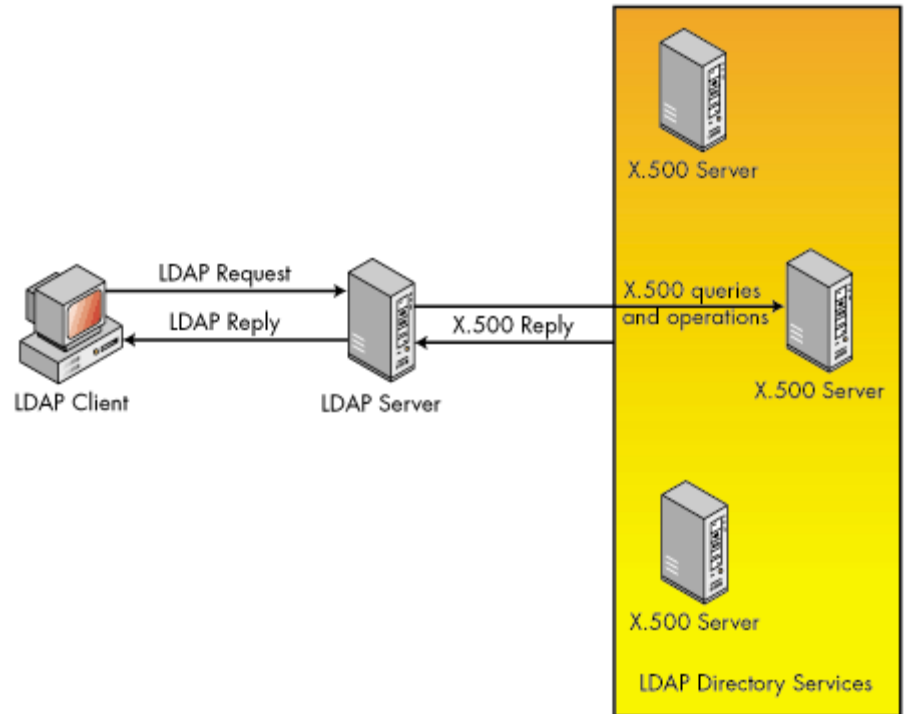
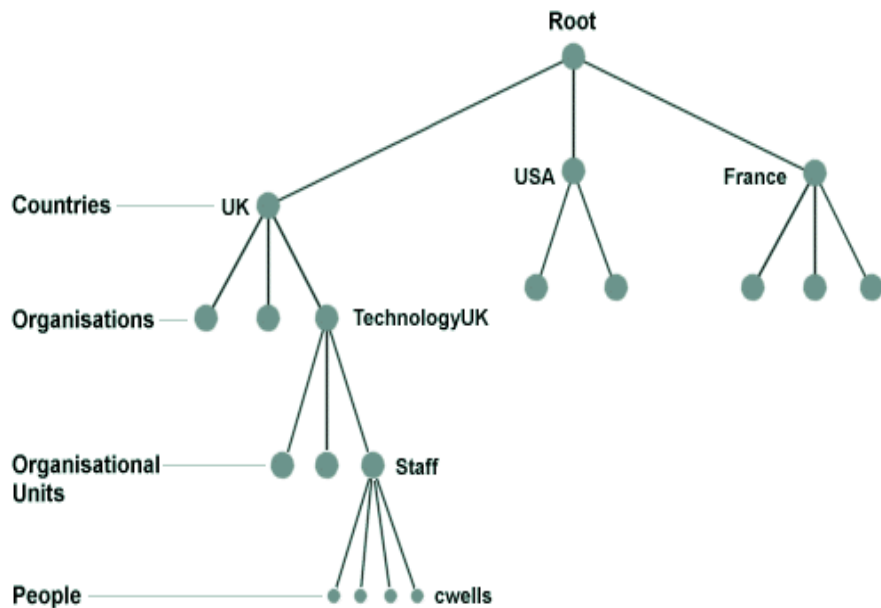
- Other name resolution mechanisms
 - multicast navigation
 - ◆ multicast name resolution request and corresponding server returns the result
 - integrated name service in V systems
 - ◆ name service is integrated with the object operation
 - server's address is returned with result of operation

Directory and Discovery Services

- Directory service
 - store collections of bindings between names and attributes
 - look up entries that match *attribute-based* specifications
- Directory service vs. name service
 - yellow page vs. white page
- Advantages of attributed-based services
 - select objects according to precise attribute specifications where names might not be known
 - hide the structure of organizations to the outside
- Discovery service
 - look up service without prior binding and user intervention for an environment where clients and services dynamically change (for mobile networks)
 - services
 - ◆ an interface automatically registering and de-registering services
 - ◆ an interface for clients to look up the services that they require from those that are currently available

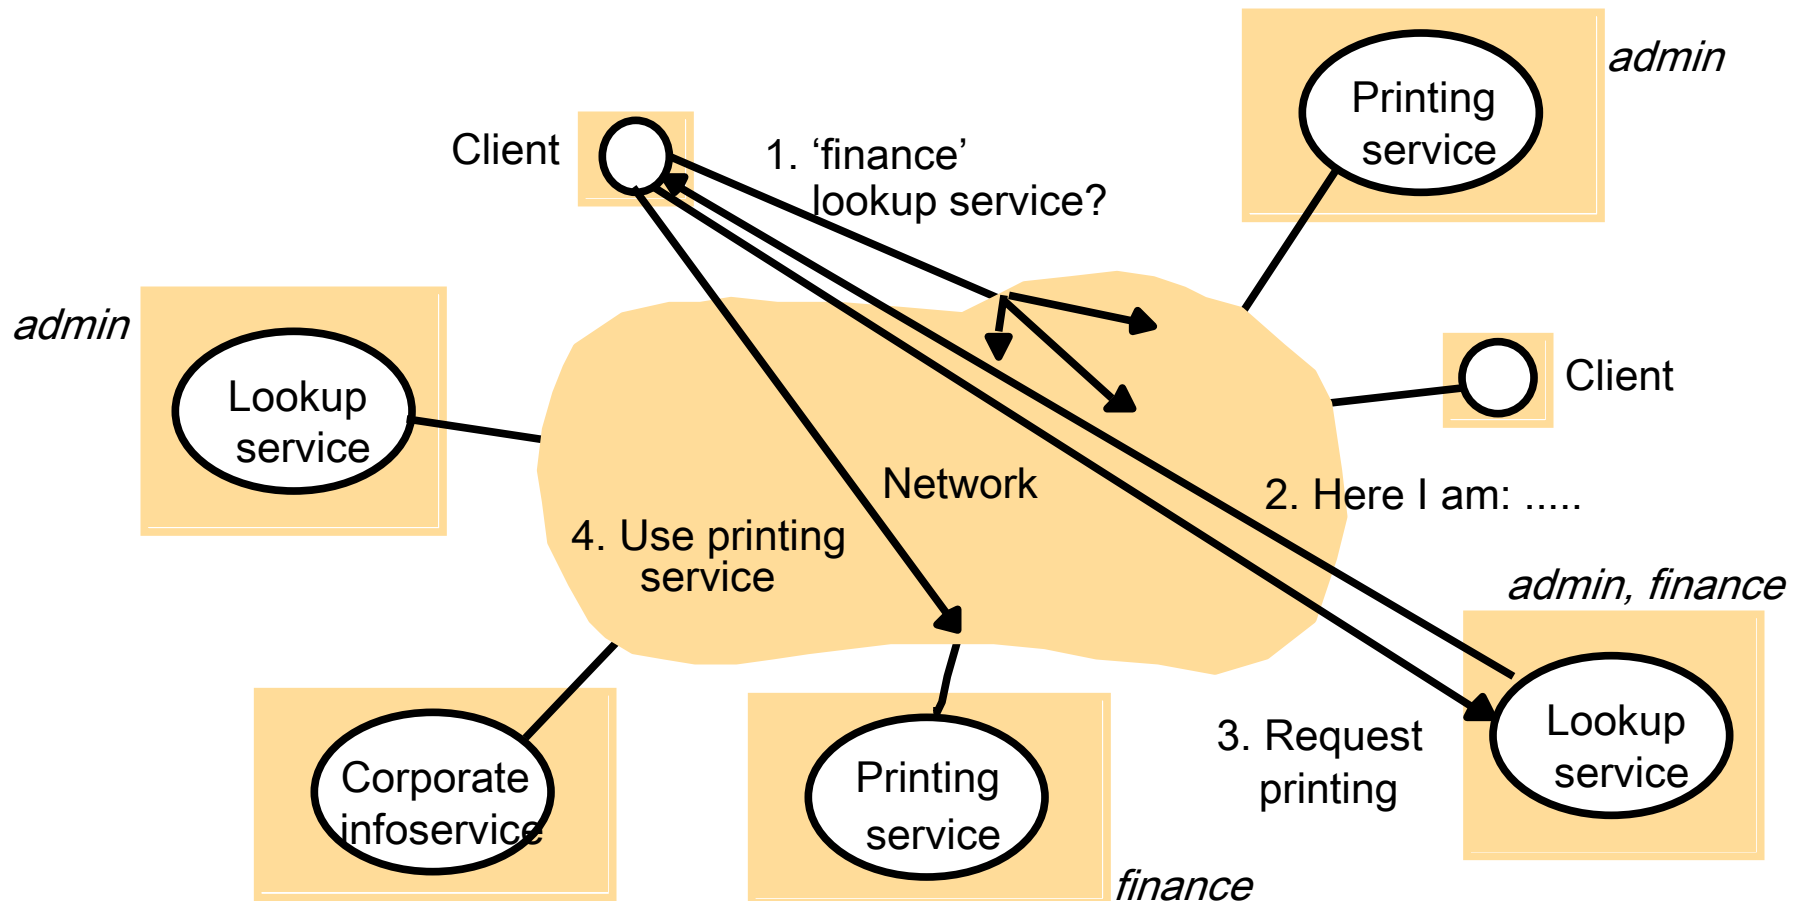
Directory and Discovery Services (cont.)

- X.500 Directory Service



Directory and Discovery Services (cont.)

- Service discovery in Jini




Directory and Discovery Services (cont.)

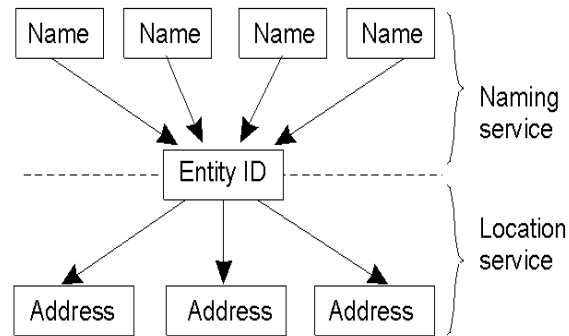
- Service discovery in ad-hoc environments
 - No central directory service
 - ♦ Jini, UPnP assume fairly stable network
 - ♦ Flooding? => designed for small network
 - Limited use of network/device resources
 - Distributed
 - ♦ Node acts as client + server
 - Service announcement
 - ♦ Flooding: How to reduce flooding traffic?
 - Periodically, delta announcement (Konark)
 - Slotted (DEAPSpace uses slotted+periodic)
 - ♦ Cache advertisements
 - ♦ TTL

Locating Mobile Entities

- Location service
 - separation between naming and locating entities via identifiers

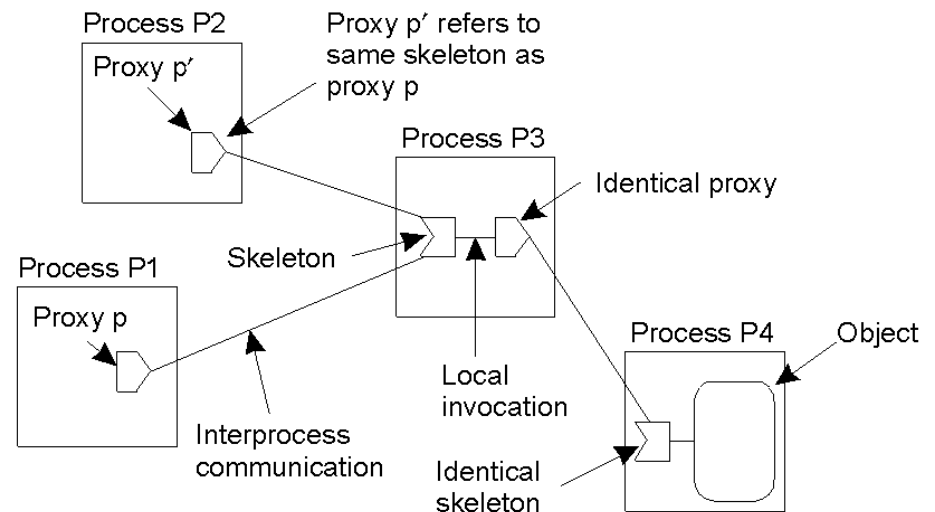


(a)



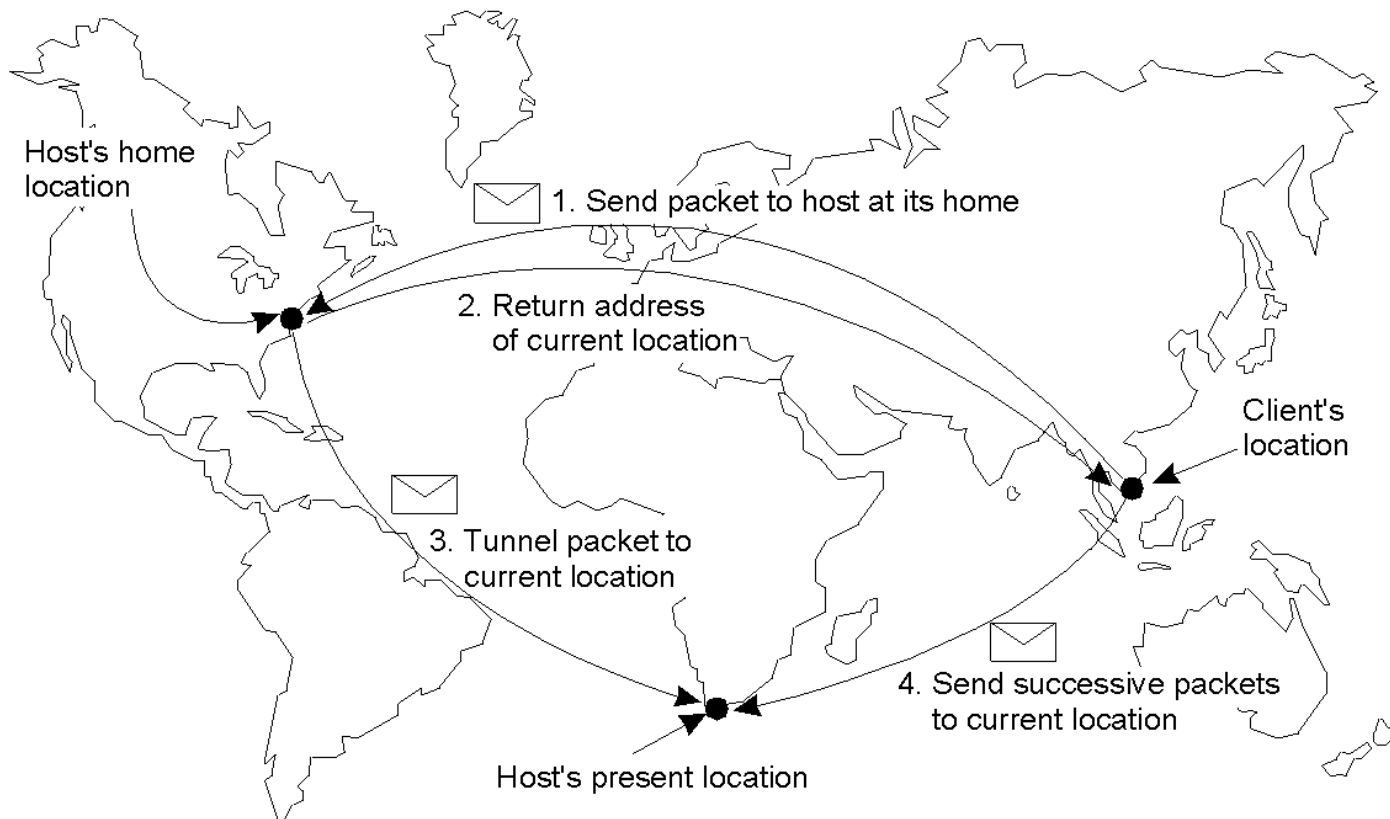
(b)

- Simple location services
 - multicasting/broadcasting
 - forwarding pointers



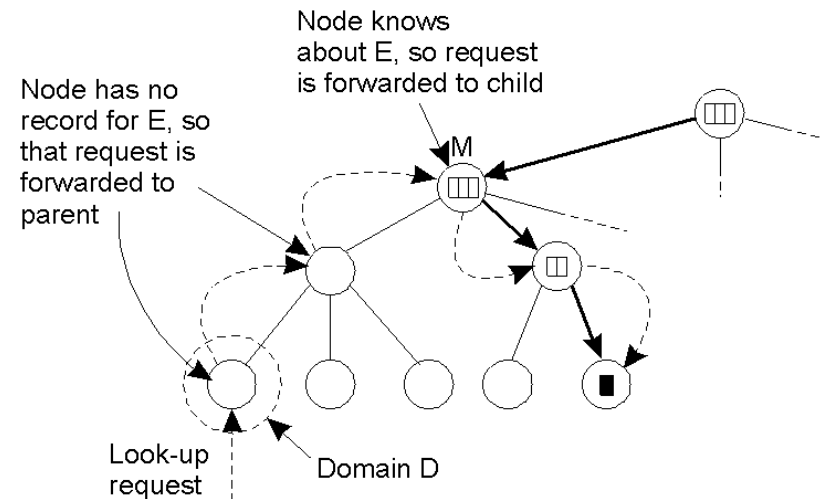
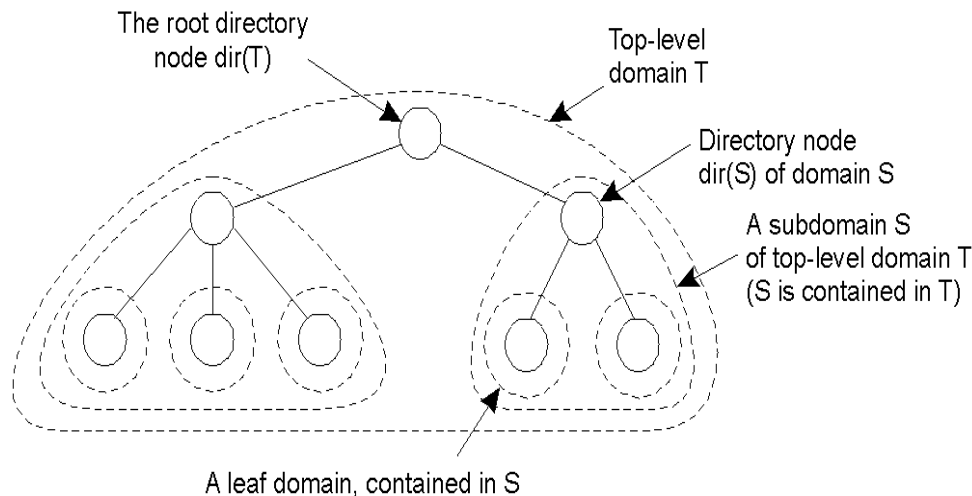
Locating Mobile Entities (cont.)

- Scalable solutions
 - home-based approach – similar to mobile IP



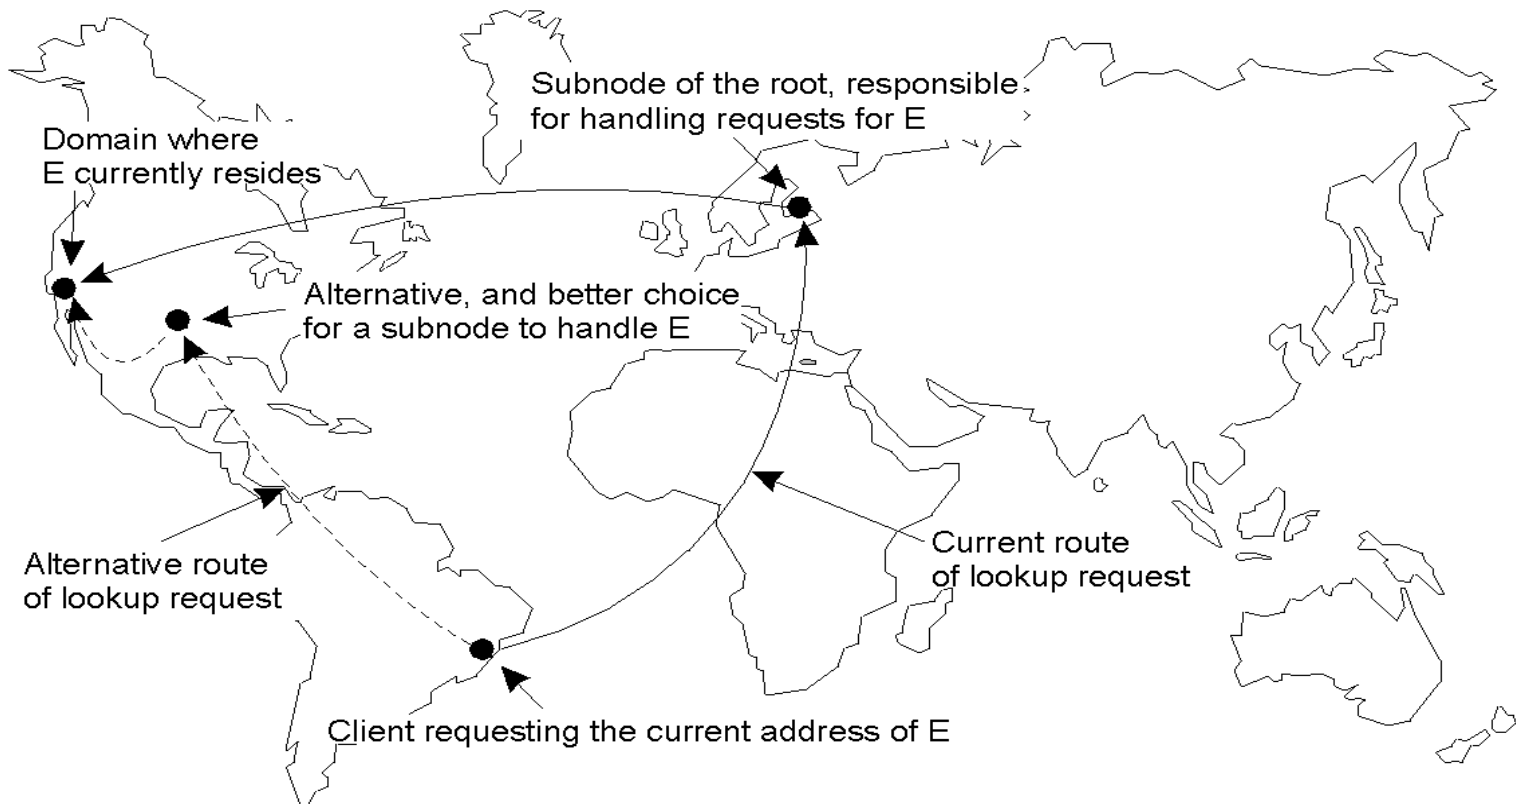
Locating Mobile Entities (cont.)

- Scalable solutions
 - hierarchical approach
 - ◆ hierarchical organization of a location service into domains, each having an associated directory node
 - ◆ pointer cache - caching a reference to a directory node of the lowest-level domain in which an entity will reside most of the time



Locating Mobile Entities (cont.)

- The scalability issues are related to uniformly placing the sub-nodes of a partitioned root node across the network covered by a location service.



Case Study: DNS

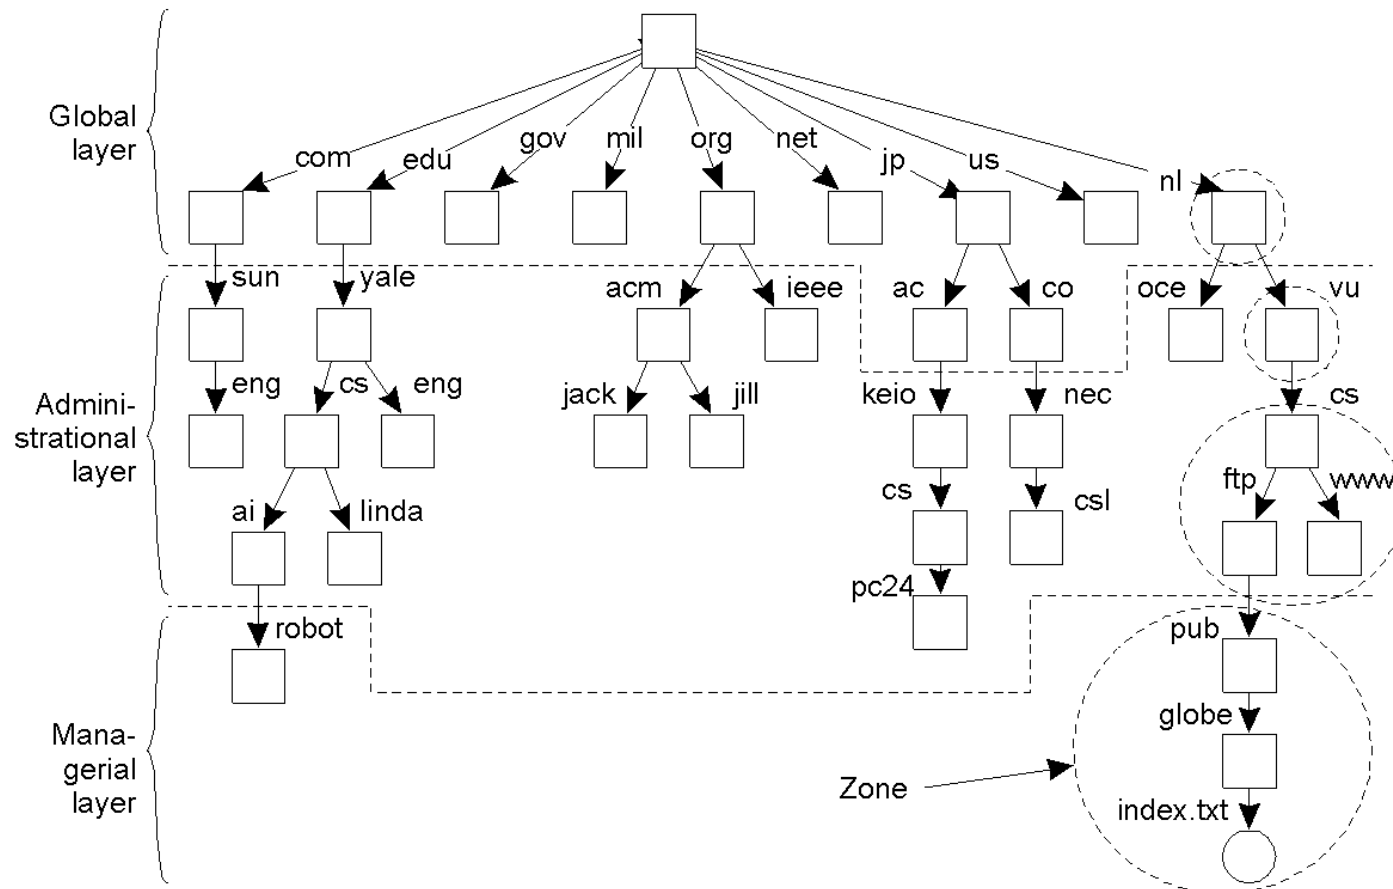
- Design goal
 - primarily used for host name resolution and for looking up mail hosts
- Characteristics
 - designed to provide short average response times for lookups
 - ◆ simple request and response protocol based on UDP
 - ◆ data are managed by a combination of partition, replication, and cache
 - temporal inconsistency is allowed
 - central authority of data consistency
 - common denominator for Internet application-level naming spaces
 - not designed for federation

Case Study: DNS (cont.)

- Name space [RFC 1034]
 - structure
 - ◆ hierarchical with a single root
 - ◆ each node has an exactly one incoming edge except the root
 - ◆ partitioned by organizations and locations
 - naming convention
 - ◆ each node is labeled with case-insensitive alpha-numeric characters
 - each label: 63 characters
 - complete path: 255 characters
 - ◆ labels are separated by “.”
 - eg: vega.kaist.ac.kr
 - ◆ read left (leaf node) to right (root or root of subtree)
 - domain name
 - ◆ list of labels on the path from node to root of sub-tree
 - ◆ relative or absolute

Case Study: DNS (cont.)

- Name space (cont.)



Case Study: DNS (cont.)

- Name service
 - DNS database is distributed across a logical network of servers
 - DNS naming data are divided into zones
 - ♦ attribute data for names in a domain
 - ♦ names and addresses of name servers
 - ♦ names of name servers for delegated sub-zones
 - ♦ zone management parameters
 - server may have data for zero or more zones
 - ♦ zone must be replicated authoritatively in at least two failure-independent servers
 - ♦ server may cache non-authoritative zone files
 - validity of cache is determined by time-to-live value

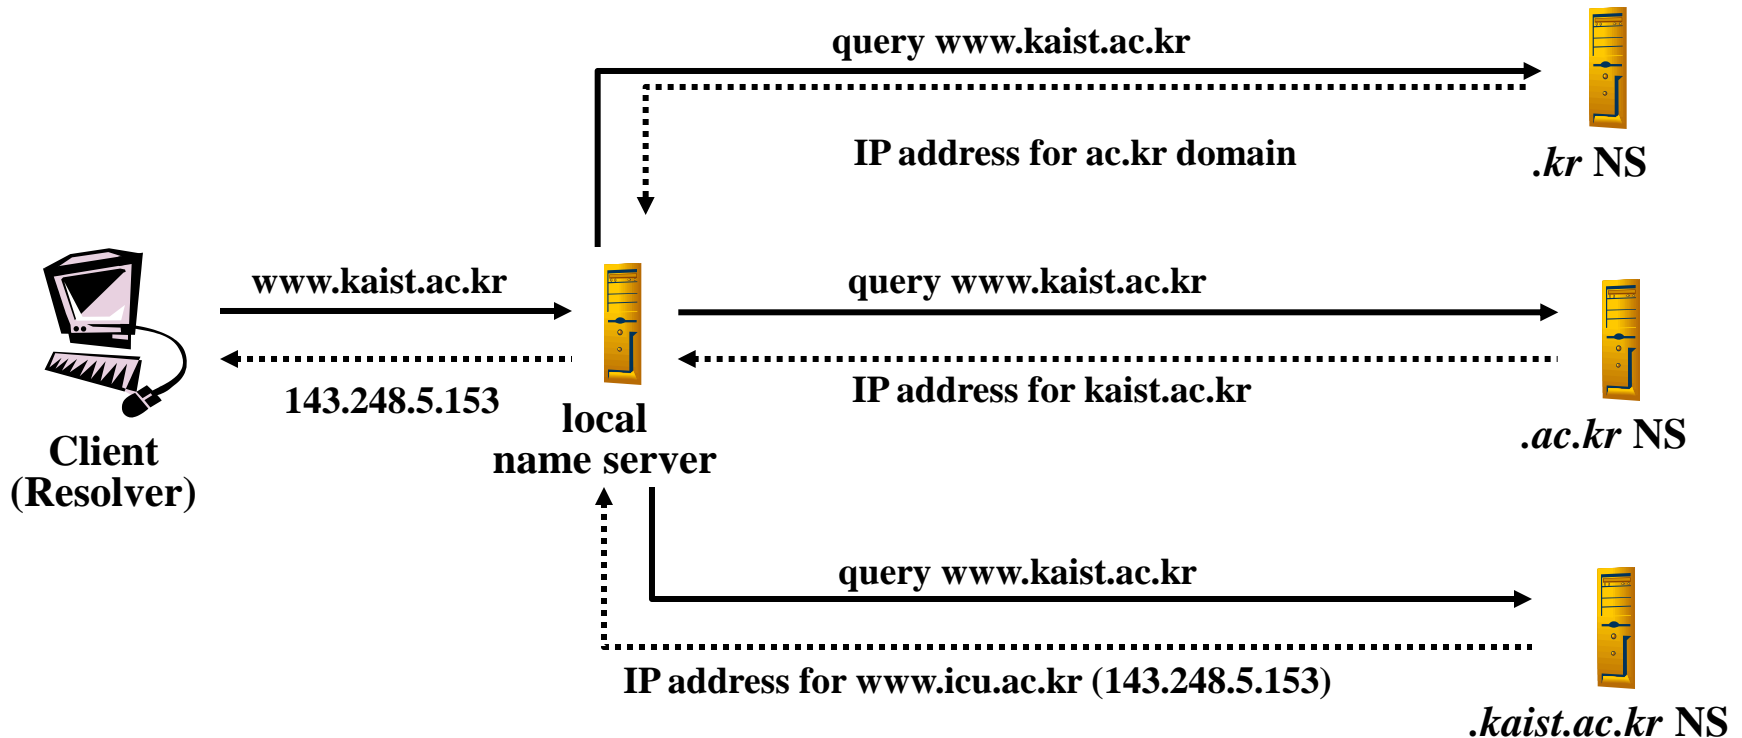
Case Study: DNS (cont.)

- Name service (cont.)
 - zone data are stored in one of several fixed types of resource records

Type of record	Associated entity	Description
SOA	Zone	Holds information on the represented zone
A	Host	Contains an IP address of the host this node represents
MX	Domain	Refers to a mail server to handle mail addressed to this node
SRV	Domain	Refers to a server handling a specific service
NS	Zone	Refers to a name server that implements the represented zone
CNAME	Node	Symbolic link with the primary name of the represented node
PTR	Host	Contains the canonical name of a host
HINFO	Host	Holds information on the host this node represents
TXT	Any kind	Contains any entity-specific information considered useful

Case Study: DNS (cont.)

- DNS name resolution [RFC 1034]
 - use DNS protocol based on UDP/IP
 - resolve chooses recursive or iterative navigation based on server's capability
 - multiple queries can be packed into a single request message



Case Study: Naming in ICN*

- ICN (Information Centric Networking)
 - Shift from host-oriented comm to content-centric
 - ◆ location-independent naming, in-network caching, and name-based routing for effective distribution of content over the network
 - general infrastructure that provides in-network caching so that content is distributed in a scalable, cost-efficient & secure manner
 - Receiver-driven model: subscribe/get objects of interest
 - Support for location transparency, mobility & intermittent connectivity
 - Needs also to be able to support interactivity (e.g. voice) and node oriented services (e.g. telnet)
- Why new naming?
 - P2P overlay and CDN are focused on how to find a local copy and do not exploit the underlying network topology

• M.F. Bari et al., “A Survey of Naming and Routing in Information-Centric Networks,” IEEE Communications Magazine, Dec 2012;
• G. Pavlou, Information-Centric Networking: Introduction and Key Issues, 2015

Case Study: Naming in ICN

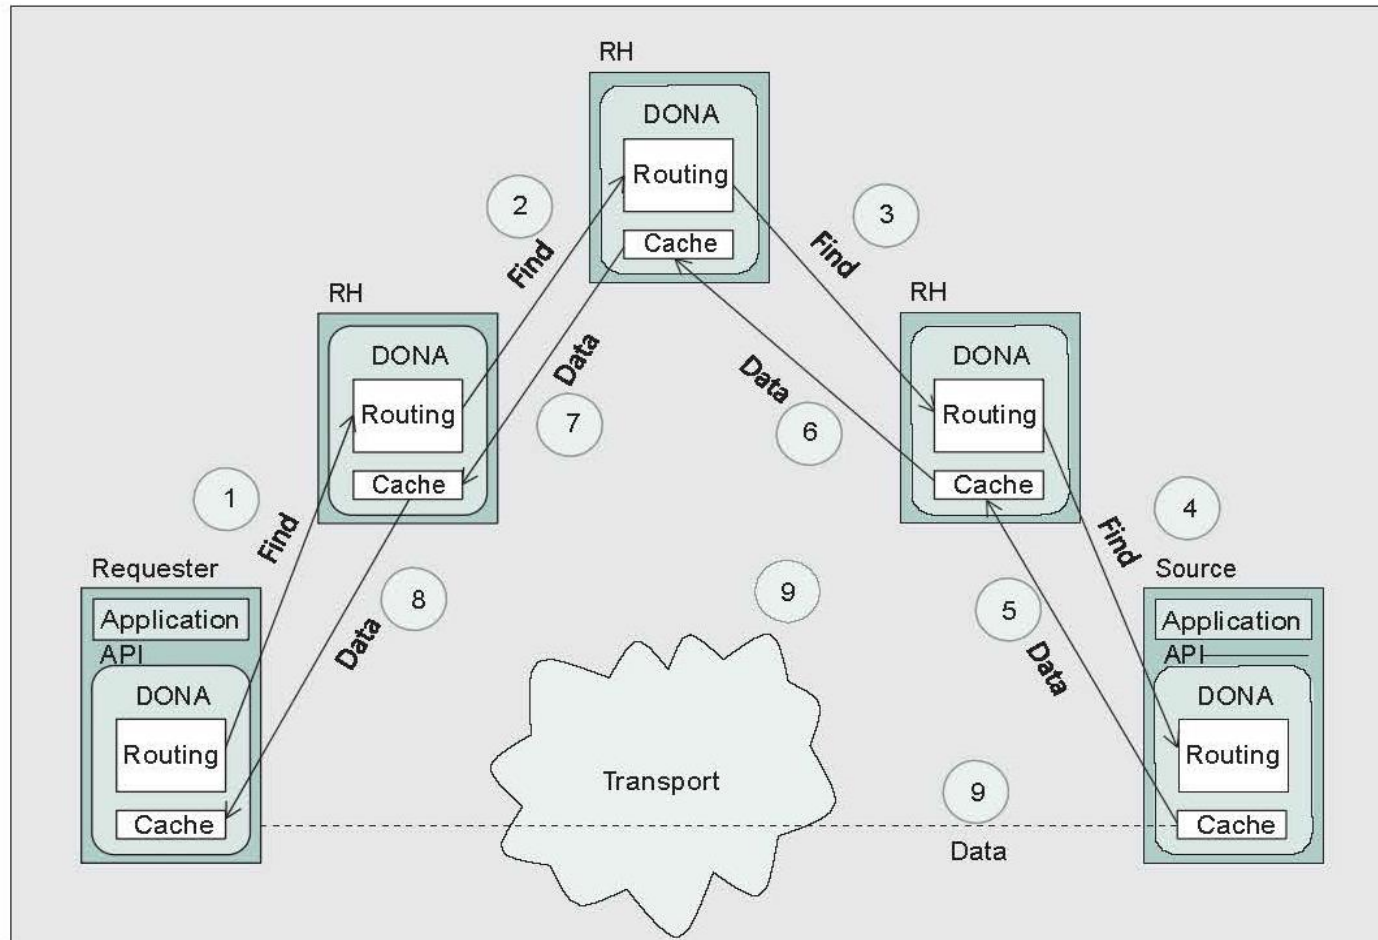
- Content Naming Issues
 - Information objects are identified by location-independent IDs, with all the object copies sharing a unique ID
 - Given that in ICN security applies to information, object IDs in many ICN architectures incorporate security
 - ◆ Non human-friendly IDs
 - ◆ Human-friendly names may additionally be associated with IDs
 - ◆ Search closely integrated with ICN given object metadata
 - Flat, hierarchical or combined ID schemes
 - Currently more than a trillion URLs (10^{12}), many more IDs are expected
 - Scalability is a concern in particular for flat naming schemes
 - ◆ Sufficient aggregation required for hierarchical schemes

Case Study: Naming in ICN

- Name Resolution and Routing Issues
 - Two approaches: dependent on namespace/ID properties
 - ◆ Two-phase: name resolution with mapping of ID to locator, and routing to the source
 - ◆ One-phase : direct ID-based routing to source
 - The two-phase approach relies on name resolution servers
 - The locator is typically not visible to the application which uses a Get(ID) API abstraction – with Put(ID) for publishing content
 - Different characteristics of the two approaches:
 - ◆ The two-phase approach can be incrementally deployed on the current Internet given that locator-based routing is used
 - ◆ The one-phase ID-based routing is radical

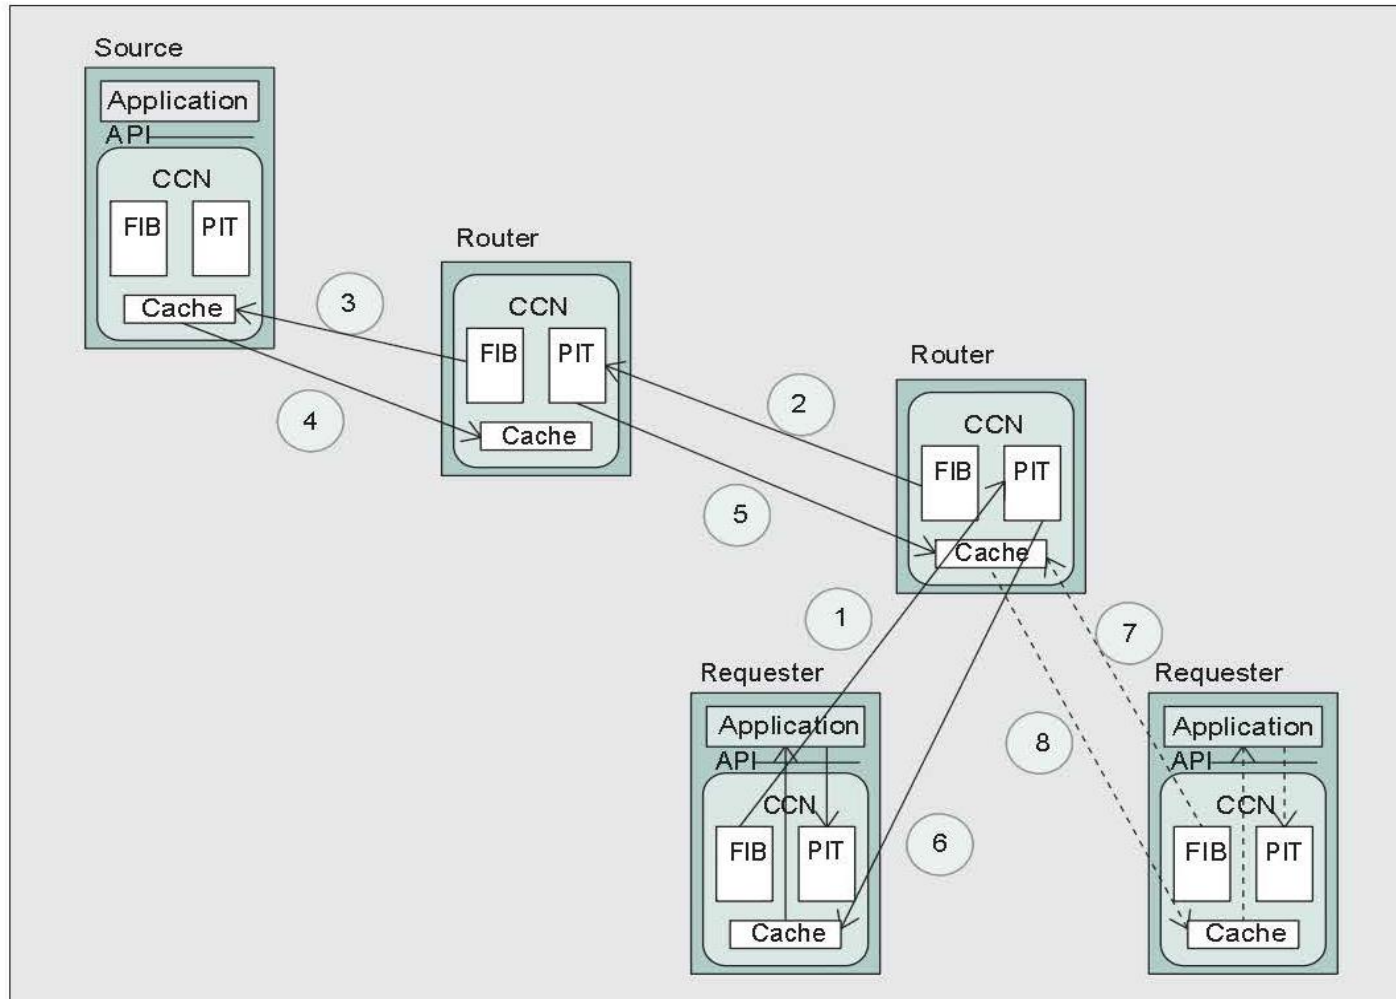
Case Study: Naming in ICN

- DONA - UCB



Case Study: Naming in ICN

- CCN – XEROX/UCLA



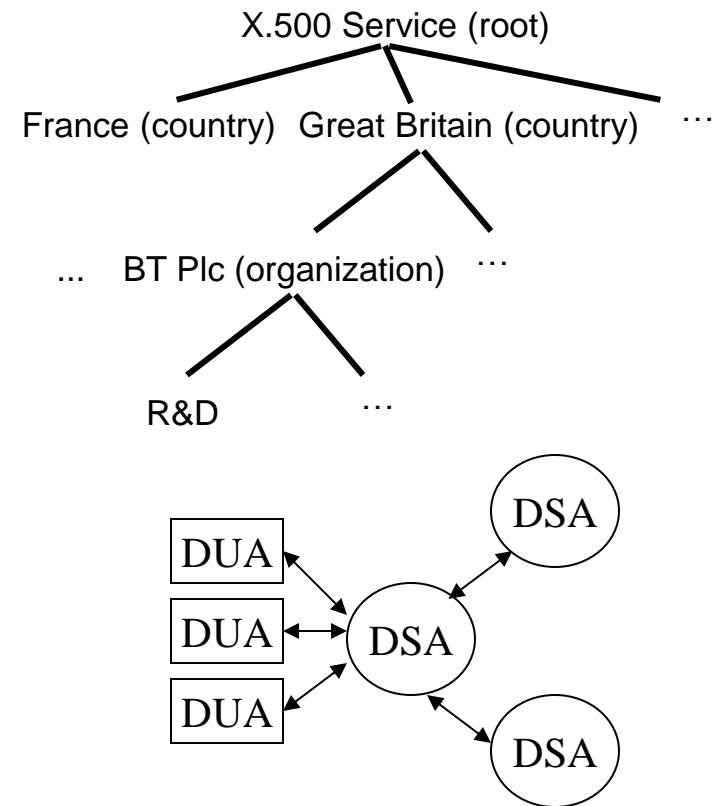
Case Study: Naming in ICN

- Name Resolution and Routing Issues
 - Two approaches: dependent on namespace/ID properties
 - ◆ Two-phase: name resolution with mapping of ID to locator, and routing to the source
 - ◆ One-phase : direct ID-based routing to source
 - The two-phase approach relies on name resolution servers
 - The locator is typically not visible to the application which uses a Get(ID) API abstraction – with Put(ID) for publishing content
 - Different characteristics of the two approaches:
 - ◆ The two-phase approach can be incrementally deployed on the current Internet given that locator-based routing is used
 - ◆ The one-phase ID-based routing is radical

Case Study: X.500

- Design goal
 - naming service supports both white page service and yellow page service
 - look up by name as well as by any required combination of attributes

- Name space
 - tree structure
 - DIT (Directory Information Tree)
 - ◆ name tree
 - DIB (Directory Information Base)
 - ◆ directory structure and attributes
 - ◆ entity: name, a set of attributes



Case Study: X.500 (cont.)

- DIB Entry

info

Alice Flintstone, Departmental Staff, Department of Computer Science,
University of Gormenghast, GB

commonName

Alice.L.Flintstone
Alice.Flintstone
Alice Flintstone
A. Flintstone

uid

alf

mail

alf@dc.s.gormenghast.ac.uk

surname

Flintstone

Alice.Flintstone@dc.s.gormenghast.ac.uk

roomNumber

Z42

telephoneNumber

+44 986 33 4604

userClass

Research Fellow

Case Study: X.500 (cont.)

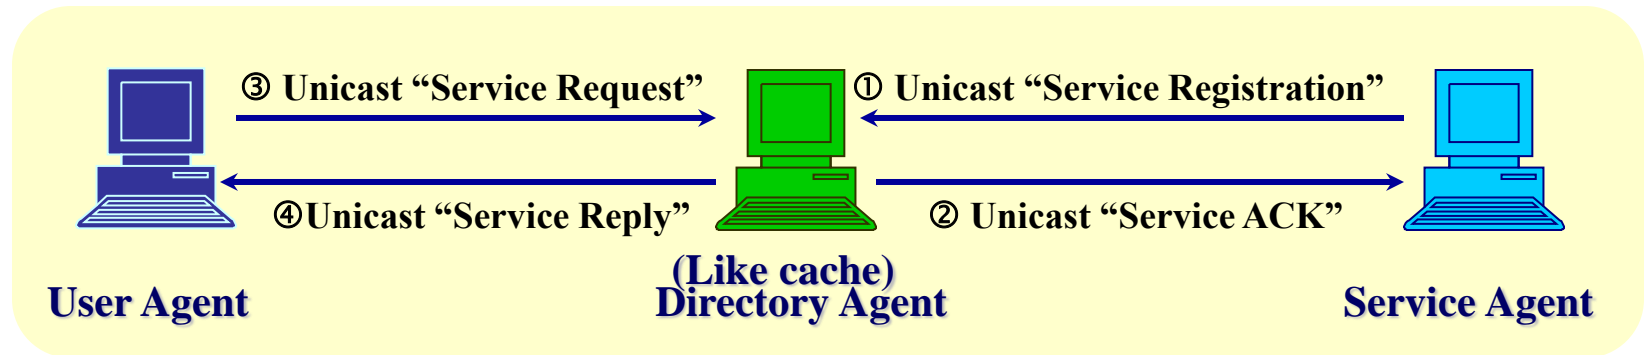
- Name service
 - read
 - ◆ DSA locates the named entry by navigating in the DIT for absolute or relative name with a list of attributes to be read and returns the required attributes
 - search
 - ◆ attribute-based access with base name and filter expression
 - ◆ return a list of names for all of entries below base node for which the filter evaluates to true
- LDAP
 - simplified interface to X.500 [RFC 2251]
 - not limited to X. 500
- Issues in X.500
 - requirement for a global directory service to exist is not clear
 - integration with Internet naming standards such as DNS
 - no clear definition for scope of information at national and international level

Case Study: SLP

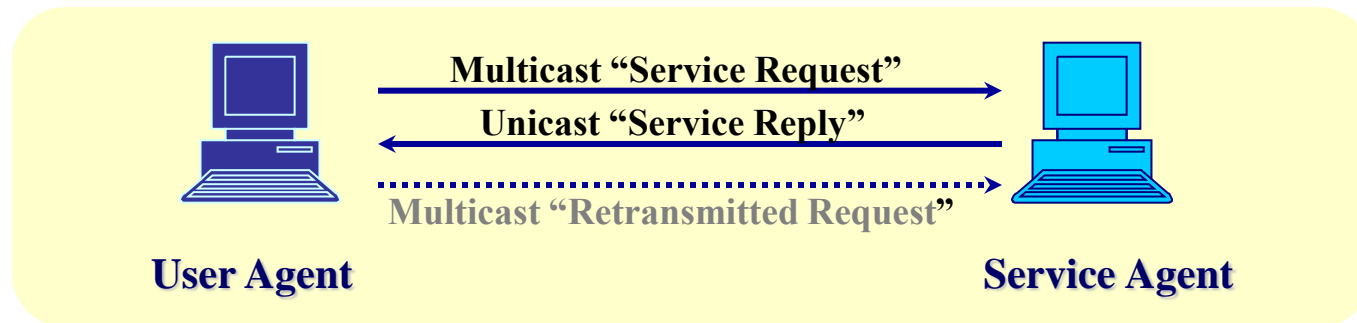
- Overview
 - IETF standard protocol for service discovery and advertisement
 - Designed solely for IP-based networks
 - Doesn't define the protocol used between the client and server
- Components
 - User Agents perform service discovery on behalf of client software
 - Service Agents advertise the location and attributes of services
 - Directory Agents aggregate service information
- Service advertisement
 - Services are advertised using service URL : IP address, port number and depending on the service type, path
 - Service templates : description of the service attributes and service scheme associated with a particular service type

Case Study: SLP – Service Discovery

- Standard case



- No DA case



Case Study: Service Discovery in Ad hoc Environments

- Service discovery schemes have been extended
 - To support P2P/Virtual backbone based operation
 - To work efficiently with ad hoc routing
- Related research
 - P2P cache : E.g. [DEAPspace]
 - P2P cache + Group : E.g. [Allia]
 - Virtual backbone with selected nodes : E.g. [ULAS03]
 - Combined with routing : E.g. [M-ZRP]