

# Distributed Systems (CS 543)

## *Replication & Fault Tolerance*

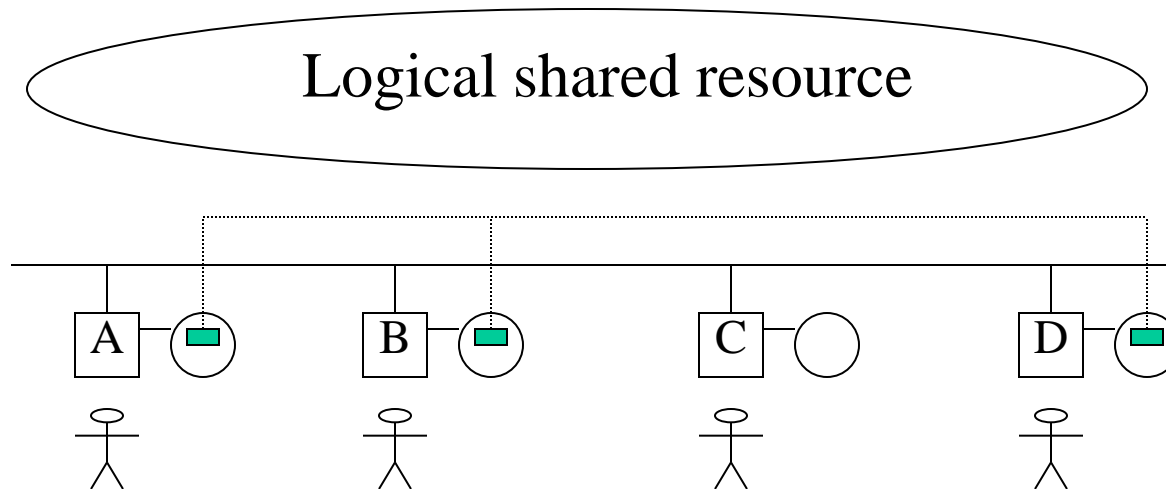
Dongman Lee  
Dept of CS  
KAIST

# Class Overview

- Why Replication?
- Replication Model
- Replication Consistency Protocols
- What is Fault?
- Fault Model
- Fault Tolerant Approaches
  - state machine
  - primary-backup

# Why Replication?

- Purpose
  - increase availability, dependability and/or performance without knowledge of replica visibility
- Replication transparency
  - hiding the replication of state in a system
    - ♦ active vs. primary/stand-by replicas
    - ♦ generic functions: active and passive replication mechanisms



# Replication Model

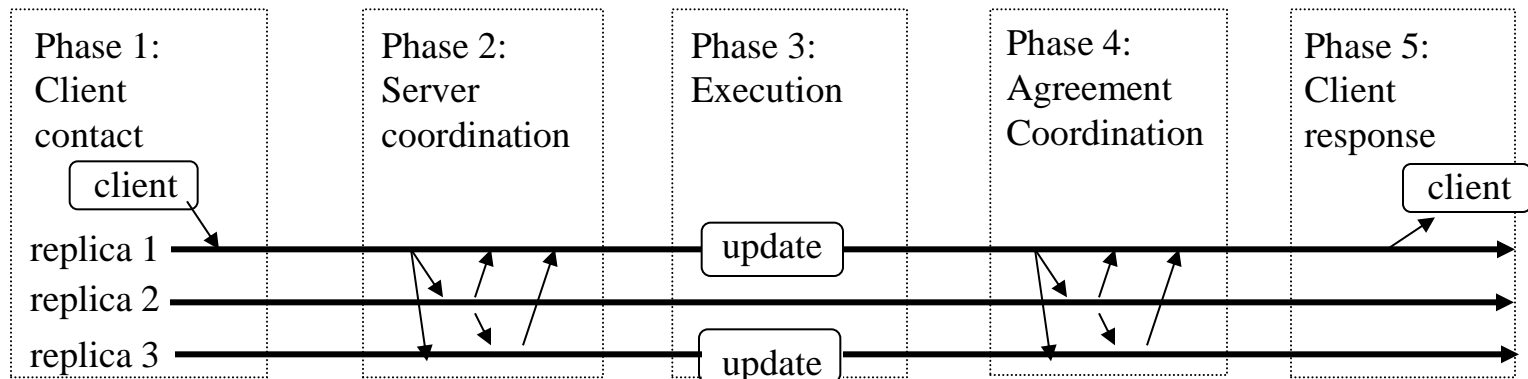
- Replication model spectrum
  - consistency
    - ◆ totally synchronous model
      - complete synchronization among replicas
    - ◆ asynchronous model
      - asynchronous update of replicas - that is, allow temporal inconsistency among replicas
    - ◆ most replication models are somewhere between these two models
  - purpose
    - ◆ performance improvement
      - reduction of delay by caching or replicating a server near clients
    - ◆ availability
      - make the service accessible (close to 100%) in the presence of process and network failures (partition and disconnection)
    - ◆ fault tolerance
      - guarantee strictly correct behavior despite of failures (byzantine and crash)

# Replication System Model (cont.)

- Replication model
  - Active replication
    - ◆ deterministic execution
    - ◆ request sent to replicas using atomic totally ordered multicast
    - ◆ no need of agreement
  - Passive replication
    - ◆ non-deterministic execution
    - ◆ view synchronization
    - ◆ no need of server coordination
  - Semi-active replication
    - ◆ non-deterministic execution
    - ◆ request sent to replicas using atomic totally ordered multicast
    - ◆ leader informs followers of its choice using view synchronization
  - Semi-passive replication
    - ◆ same as passive without view synchronization
    - ◆ allow for aggressive time-outs values and suspecting crashed processes without incurring too high cost for incorrect failure suspicions

# Replication System Model [Weismann]

- Replication protocol model
  - Request phase
    - ◆ active replication
    - ◆ passive replication
  - Server coordination
    - ◆ message ordering: FIFO, causal, total
  - Execution
  - Agreement coordination
    - ◆ necessary in database while ordering guarantee is enough for distributed systems
  - Client response
    - ◆ synchronous vs. lazy or asynchronous




# Replication Consistency Protocols

- Description
  - describe an implementation of a specific consistency model
- Classification
  - primary-based protocols
    - ◆ remote-write protocols
    - ◆ local-write protocols
  - replicated-write protocols
    - ◆ active replication
    - ◆ quorum-based protocols

# Primary-based Remote-Write Protocols

- All write operations are performed at a (remote) fixed server
  - read operations are allowed on a local copy while write operations are forwarded to a fixed primary copy



W1. Write request  
W2. Forward request to primary  
W3. Tell backups to update  
W4. Acknowledge update  
W5. Acknowledge write completed

R1. Read request  
R2. Response to read

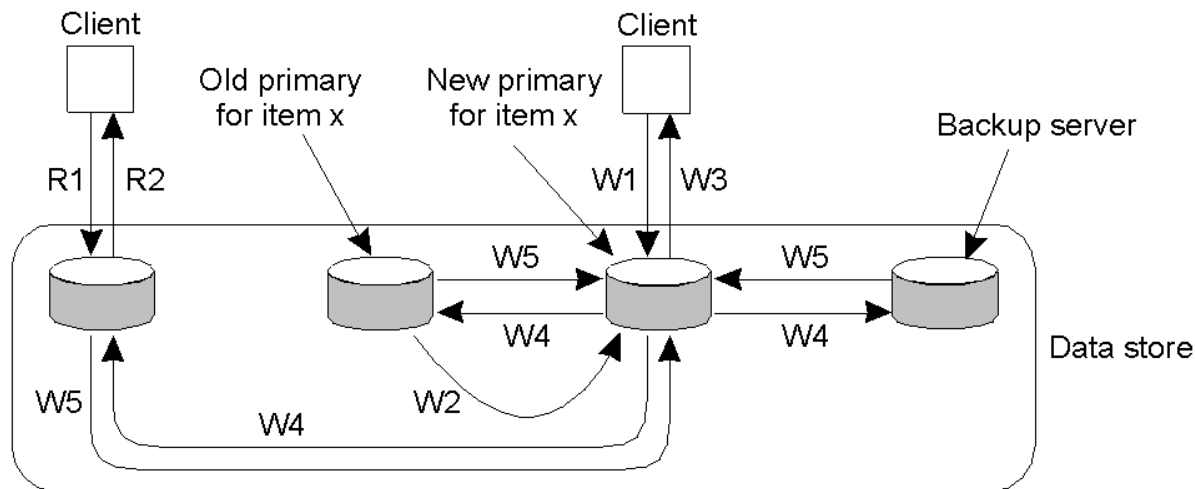


# Primary-based Remote-Write Protocols (cont.)

- Issues
  - update can be a performance bottleneck if implemented as a blocking operation
    - ◆ but guarantees sequential consistency (most recent write as the result of a read)
    - ◆ if implemented as a non-blocking, the protocol provides no guarantee of sequential consistency and fault tolerance

# Primary-based Local-Write Protocols

- All write operations are performed locally and forwarded to the rest of replicas
  - primary copy migrates between processes that wish to perform a write operation
  - Multiple, *successive* writes can be done locally (via non-blocking protocol)
  - can be exploited in mobile computing




W1. Write request  
W2. Move item x to new primary  
W3. Acknowledge write completed  
W4. Tell backups to update  
W5. Acknowledge update

R1. Read request  
R2. Response to read


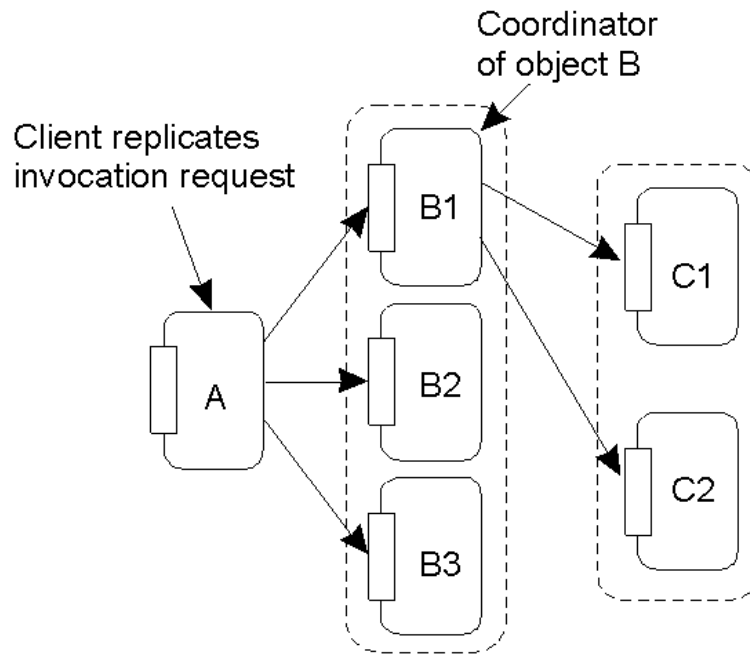
# Active Replication

- Each replica performs update operations and propagates them (or the results) to the others
  - requires totally ordered multicast
- Replicated invocation problem



# Active Replication (cont.)

- Solutions to the replicated invocation problem
  - group coordinator
  - sender-driven vs. receiver-driven




# Quorum-based Protocols

- Require clients to request and acquire the permission of multiple servers before any operation on replicas
  - quorum set
    - ◆  $W > \text{half the total votes}$
    - ◆  $R + W > \text{total number of votes for group}$ 
      - any pair of read quorum and write quorum must contain common copies, so no conflicting operations on the same copy
    - ◆ read operations
      - check if there is enough number of copies  $\geq R$
      - perform operation on up-to-date copy
    - ◆ write operations
      - check if there is enough number of up-to-date copies  $\geq W$
      - perform operation on all replicas

# Quorum-based Protocols (cont.)

- Examples



- a) A correct choice of read and write set
- b) A choice that may lead to write-write conflicts since  $W \leq N/2$
- c) A correct choice, known as ROWA (read one, write all)


# What is Fault?

- Definition
  - system is considered *faulty* once its behavior is no longer consistent with its specification [Schneider]
- Separation property of distribution systems lead to *partial failure property*
  - components that one component depends on may fail to respond due to various reasons
    - ◆ system or network failure
    - ◆ system or network overload
- Dependability
  - availability: readiness of use
  - reliability: continuity of service delivery
  - safety: low chances of catastrophes
  - maintainability: repairability

# Failure Model

- Failure semantics

- description of the ways in which a service may fail



- recovery actions depend on the likely failure behavior of a server when its failure is detected
- designers should ensure that the behavior of a server conforms to a specified failure semantics
  - ◆ e.g. network with omission/time failure semantics
    - need to guarantee detection of message corruption such as checksum
  - ◆ stronger failure semantics costs more in general
- adequacy of failure semantics would require preliminary stochastic analyses



# Failure Model (cont.)

- Representative faulty behavior
  - Fail-stop failures
    - ◆ when system fails, it changes to a state that allows others to detect its failure and then stops
  - Byzantine failures
    - ◆ system exhibits arbitrary and malicious behavior which may collude with other systems

# Fault-Tolerant Approaches

- Fault tolerance
  - can detect a fault and either fail predictably or mask the fault from users
  - hiding the occurrence of errors in system components and communications
  - ➔ incorporate *redundant* processing component to achieve fault tolerance
- *k-resilient/fault-tolerant*
  - a set of systems satisfies its specification if no more than  $k$  systems become faulty
  - $k$  is chosen based on statistical measures of system reliability
    - ◆ Arbitrary failure:  $2k+1$  (identical group);  $3k+1$  (non-identical; Byzantine failure)
    - ◆ fail-stop failure:  $k+1$

# Fault-Tolerant Approaches (cont.)

- Two approaches to support fault tolerance (fault masking)
  - hierarchical failure masking
    - ◆ hierarchical failure and recovery management
      - error detection in layered communication protocols
      - various levels of error abstraction in OS
  - group failure masking
    - ◆ state-machine approach
    - ◆ primary-backup approach
- Fault tolerance support can be done
  - hardware
    - ◆ stable storage
  - software
    - ◆ replicated servers

# State-Machine Approach

- Requirements for  $k$  fault-tolerant state machine
  - all replicas receive and process the same sequence of requests
    - ◆ agreement: every non-faulty replica receives every request
      - specify the interaction behavior of a client with state machine replicas
      - relaxed for read-only request in fail-stop failures
    - ◆ order: every non-faulty replica processes requests it receives in the same relative order
      - specify the behavior of state machine replicas in term of how to process requests from clients
      - relaxed for commutative requests

# State-Machine Approach (cont.)

- Agreement requirement
  - to satisfy agreement requirement, state-machines should support a message broadcasting protocol which conforms to
    - ♦ IC1: all non-faulty processors agree on the same value
    - ♦ IC2: if sender of request is non-faulty, then all non-faulty processors use its value as the one on which they agree
  - message broadcasting protocol is called Byzantine agreement protocol or reliable broadcast protocol

# State-Machine Approach (cont.)

- Order requirement
  - to implement order requirement requires
    - ◆ assignment of unique identifier to each message
    - ◆ stability (a request is ready to be delivered once all the previous requests have been delivered) test
  - assumptions on order requirement
    - ◆ O1: requests issued by a single client to a given state machine  $sm$  are processed by  $sm$  in the order they were issued
    - ◆ O2: if the fact that a request  $r$  was made to a state machine  $sm$  by a client  $c$  could have caused a request  $r'$  to be made by a client  $c'$  to  $sm$ , then  $sm$  processes  $r$  before  $r'$
  - three approaches
    - ◆ logical clock-based
    - ◆ synchronized real-time clock-based
    - ◆ replica-generated identifiers-based

# State-Machine Approach (cont.)

- Order requirement: logical clock-based
  - only for failstop failures
  - unique id assignment: logical clock
    - ♦ LC1: timestamp is incremented after each event at  $p$
    - ♦ LC2: upon receipt of a message with timestamp  $t$ , process  $p$  resets its timestamp  $T_p$  to  $\max(T_p, t)+1$
  - stability test
    - ♦ a request is stable at replica  $sm_i$  if a request with larger timestamp has been received by  $sm_i$  from every client running on a non-faulty processor
      - messages between a pair of processors are delivered in the order sent
      - processor  $p$  detects that a failstop process  $q$  has failed only after  $p$  has received  $q$ 's last message sent to  $p$

# State-Machine Approach (cont.)

- Order requirement: synchronized physical clock-based
  - unique id assignment
    - ♦ no client makes two or more requests between successive clock ticks  
=> every message will have greater timestamp than its previous message (satisfies O1)
    - ♦ degree of clock synchronization is better than minimum message delivery time => timestamps of two causally related messages issued by two clients will be such that earlier one should have lower timestamp than later one (satisfies O2)
  - stability test
    - ♦ request  $r$  is *stable* if local clock reads  $T$  and  $uid(r) < T-d$  ( $d$ : worst case message delivery time)
    - ♦ request  $r$  is *stable* if a request with larger uid has been received from every client



# State-Machine Approach (cont.)

- Order requirement: replica-generated identifiers-based
  - 2 phases are used
    - ♦ phase 1: replicas propose uid as part of agreement protocol (SEEN)
    - ♦ phase 2: one of candidates is selected and becomes uid (ACCEPTED)
  - stability test
    - ♦ request  $r$  that has been accepted by  $sm_i$  is *stable* if there is no request that has
      - been seen by  $sm_i$ ,
      - not been accepted by  $sm_i$ , and
      - for which  $cuid(sm_i, r) \leq uid(r)$  holdswhere  $cuid(sm_i, r) = \max(\text{SEEN}_i, \text{ACCEPT}_i) + 1 + i$   
 $\text{SEEN}_i$ : largest  $cuid(sm_i, r)$  assigned to any request  $r$  so far seen by  $sm_i$   
 $\text{ACCEPT}_i$ : largest uid( $r$ ) assigned to any request  $r$  so far accepted by  $sm_i$   
 $uid(r) = \max_{sm_j \in NF} (cuid(sm_j, r))$  where  $NF$  be the set of replicas from which candidate unique identifiers(*cuid*'s) were received

# Primary-Backup Approach

- Cost metrics of primary-backup protocols
    - degree of replication
      - ♦ # of servers for fault tolerance
    - blocking time
      - ♦ worst cast period between a request and its response in any failure-free execution
    - failover-time
      - ♦ worst-case period during which requests can be lost because there is no primary
- ⇒ Smallest degree of replication, blocking time, failover-time for k-fault-tolerance?

# Primary-Backup Approach (cont.)

- Protocol properties
  - Pb1: there is at most one server whose state satisfies a condition being a primary
    - ♦ no more than one server is the primary at a time
  - Pb2: each client maintains a server identity to which the client can send a message
    - ♦ a client sends a request to the service by sending it to the server it believes to be the primary
  - Pb3: if a client request arrives at a server that is not a primary, then that request is not enqueued (thus, not processed)
    - ♦ messages to a backup are ignored
  - Pb4: there exist fixed value  $k$  and  $\Delta$  such that the service behaves like a single  $(k, \Delta)$ -bofo server\*
    - \*  $(k, \Delta)$ -bofo server (bounded outage, finitely often) : all server failures can be grouped into at most  $k$  intervals of time with each interval having length at most  $\Delta$

# Primary-Backup Approach (cont.)


- Simple primary-backup protocol

- assumption

- ◆ one primary server  $p_1$  and one backup server  $p_2$ , connected via a communication link (message delivery time upper bound:  $\delta$ )
    - ◆ operations when  $p_1$  receives a request from a client
      - processes the request and updates its state
      - send update info to  $p_2$  (a state update message)
      - send a response to the client without waiting for ack from  $p_2$
    - ◆  $p_1$  sends a dummy message every  $\tau$  seconds; If  $p_2$  does not receive a dummy message for  $\tau + \delta$  seconds,  $p_2$  becomes a primary

- spec conformance

- ◆ Pb1:  $(p_1 \text{ has not crash}) \wedge (p_2 \text{ has not received a message from } p_1 \text{ for } \tau + \delta) = \text{false}$
    - ◆ Pb2: client  $c$  sends a message to  $p_1$
    - ◆ Pb3: requests are not sent to  $p_2$  until after  $p_1$  has failed
    - ◆ Pb4: a single  $(1, \tau + 4\delta)$ -bofo server



# Primary-Backup Approach (cont.)

- Failure models
  1. crash failures
    - ◆ permanent halt - once a server halts, it never recovers
  2. crash + link failures: 1+ link may lose messages
    - ◆ links do no delay, duplicate or corrupt messages
  3. receive-omission failures: 1+ failed to receive some of messages
  4. send-omission failures: 1 + failed to send some of replies
  5. general-omission failures: 3 + 4

# State-machine vs. Primary-backup

- Comparison

	State-machine	Primary-backup	Remarks
Arbitrary Failure support	Yes	No	$2k+1/3k+1$ replication for $k$ -resilience
Request loss	No	Possible	Loss happens when a primary fails
Failure handling	Voting	Failover	
Request copy	as many servers as $k$ -resilience suffices	Only to primary	$2k+1/3k+1$ for arbitrary $k+1$ for fail-stop
Overall cost	expensive	cheap	Primary-backup approach is more popular in commercial applications